

EquiOrient Phase 2: Supplementary Material

001

Anonymous WACV [Datasets Track](#) submission

002

Paper ID *****

1. Full 45-test verification suite

003

All 45 tests must pass before any GPU run. The suite catches the entire class of bugs discovered in Phase 1 (zero structural losses, wrong label copying, stale data volumes).

004

005

Category	Tests
D4 group axioms (4)	identity, $R^4=I$, $H^2=I$, $HRH=R^{-1}$
Label action (3)	direction_H, direction_R, label_composition_all
Renderer (4)	H/R/R2/R3 pixel transform, bbox transform, object identity
Dataset integrity (5)	scene disjoint, balanced labels, no boundary, transform coverage, unique pngs
Structural loss (4)	equivariance nonzero, wrong nonzero, wrong differs, manipulation check
Gradient flow (5)	grad reaches pair_encoder (x2), vision_lora, rho never receives label, common init
Architecture (3)	matched init, equal param count, z_ablation
Collapse checks (4)	variance, effective rank, norm, dead neurons
Identifiability (2)	collision matrix, unseen distinguishes correct/wrong

Table 1. The 45 pre-GPU verification tests.

2. Full per-transform accuracy (5-seed mean)

006

Arm	I	R	R2	R3	H	RH	R2H	R3H
Original SFT	0.999	1.000	1.000	1.000	1.000	0.999	0.999	0.999
Augmentation	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000
Output cons.	1.000	0.999	1.000	0.999	0.999	0.999	0.999	0.998
Latent inv.	0.999	1.000	1.000	1.000	1.000	1.000	1.000	1.000
EquiOrient	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000
Wrong geom.	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000

Table 2. Per-transform accuracy on test set (5-seed mean). Held-out elements: R2, R3, RH, R2H, R3H. All arms at ceiling.

3. Training loss curves (per epoch)

007

4. Individual seed results

008

5. Identifiability audit

009

The collision matrix compares $\rho(g)$ vs. $\tilde{\rho}(g)$ for all 8 group elements:

010

Arm	ans ep1	ans ep2	struct ep1	struct ep2
Augmentation	3.217	0.019	—	—
Output cons.	5.109	0.986	0.686	0.459
Latent inv.	4.337	0.606	0.303	0.222
EquiOrient	4.106	0.211	0.318	0.108
Wrong geom.	4.424	0.388	0.351	0.174

Table 3. Per-epoch answer loss and structural loss (seed 101, $N=512$). Structural losses are nonzero and convergent (manipulation check passed).

Arm	s101	s202	s303	s404	s505
Original SFT	0.9984	1.0000	0.9988	0.9998	1.0000
Augmentation	1.0000	1.0000	1.0000	1.0000	1.0000
Output cons.	0.9963	0.9996	1.0000	1.0000	1.0000
Latent inv.	0.9990	1.0000	1.0000	1.0000	1.0000
EquiOrient	1.0000	1.0000	1.0000	1.0000	1.0000
Wrong geom.	1.0000	1.0000	1.0000	1.0000	1.0000

Table 4. Unseen-group accuracy per seed.

Element	$\rho(g) = \tilde{\rho}(g)?$	Distinguishes?
I	yes (trivially)	no
R	no ($R_{90} \neq -I$)	yes
R^2	no ($-I \neq I$)	yes
R^3	no ($R_{270} \neq -I$)	yes
H	yes (both = H)	no
RH	no ($R_{90}H \neq -IH$)	yes
R^2H	no ($-H \neq H$)	yes
R^3H	no ($R_{270}H \neq -IH$)	yes

Table 5. Collision matrix. The 5 held-out elements all distinguish correct from wrong representations. No collisions on the unseen set.

011 6. Bootstrap analysis

012 Hierarchical paired bootstrap: resample training seeds, then test scene IDs within each seed, preserving method
013 pairing.

- 014 • Mean ΔA (EquiOrient – Augmentation): 0.0000
- 015 • 95% CI: [0.0000, 0.0000]
- 016 • Individual seed deltas: all exactly 0.0000
- 017 • Minimum meaningful effect: 0.03 (3pp)
- 018 • Conclusion: no evidence of meaningful behavioral benefit

019 7. Dataset statistics

020 8. Compute ledger

021 9. Reproducibility

- 022 • Code: All code on branch research/equiorient-phase2 at commit 9af8d25.
- 023 • Results: 30 JSON files in results/phase2_confirmatory/, one per (arm, seed) combination.
- 024 • Frozen ledger: equiorient/freeses/phase2_confirmatory.yaml committed before inspecting confirmatory results.
- 025 • Modal scaffold: modal/equiorient_phase2.py (serverless, pinned stack: torch 2.8.0+cu128, transformers==4.57.6,
026 Qwen3-VL-8B @ 0c351dd0).

Property	Value
Scenes (total)	4,096
Scenes (train pool)	2,048
Scenes (dev)	512
Scenes (val)	512
Scenes (test)	1,024
Examples per scene	8 (all D4 elements)
Total examples	32,768
Target size range	2–4 px
Distractors per scene	12–20
Color palette	3 muted hues
Background	(155, 152, 148)
Per-pixel noise	amplitude 12
Canvas size	192×192 px
Sparse exposure	identity + 1 generator (50/50 H/R)

Table 6. v4 dataset parameters.

Run	Platform	Outcome	GPU time
v1 dev (2 jobs)	L40S (Modal)	passed, ceiling	~34 min
v3 dev (2 jobs)	L40S (Modal)	passed, ceiling	~34 min
v4 dev (2 jobs)	L40S (Modal)	passed, ceiling	~34 min
Confirmatory (30)	L40S (Modal)	VALID	~8.5 h
Total			~9.5 h

Table 7. All GPU runs. Dev runs established the ceiling; confirmatory runs collected the final results.

- Verification: 45-test pre-GPU gate (equiorient/tests/test_phase2_gate.py) must pass before any GPU run.