

Beyond Spatial Accuracy: Diagnosing Persistent Orientation Failures in Vision–Language Models

Anonymous WACV Datasets Track submission

Paper ID *****

Abstract

001 Aggregate benchmark accuracy can conceal capability-
002 specific failures, and it cannot tell an evaluator
003 whether an apparent weakness is genuine, annotation-
004 driven, representation-limited, logically incoherent,
005 benchmark-specific, or model-family-specific. We in-
006 troduce the Capability Audit Ladder, a reusable eval-
007 uation protocol that holds a relation family and its ex-
008 amples fixed while testing a suspected failure across
009 seven stages—detection, scaling, adaptation, label re-
010 audit, representation readout, logical reformulation,
011 and cross-dataset transfer—each with a pre-specified
012 decision outcome. We instantiate it on object-relative
013 orientation in vision–language models (VLMs), a small
014 but controlled case study. Scaling and adaptation
015 substantially improve overall spatial accuracy while
016 leaving orientation comparatively resistant (a 62–66%
017 band across trained conditions and a 2026 sparse-MoE
018 model); label re-audits by two independent human
019 raters confirm the pattern is not an annotation artifact;
020 representation readouts and complementary-statement
021 probes show accuracy and logical coherence are dis-
022 tinct axes that both remain weak; and benchmark-
023 specific adaptation negatively transfers to SITE’s offi-
024 cial spatial-relation category. External zero-shot mod-
025 els bound the pattern’s universality. The protocol is
026 specified operationally, every headline number is trace-
027 able to released per-example predictions, and the same
028 ladder applied to a scaling-responsive family (horizon-
029 tal) yields a different diagnosis, demonstrating discrim-
030 inability.

031 1. Introduction

032 This paper contributes an evaluation practice as much
033 as a set of model results. We develop a protocol for
034 testing, stress-testing, auditing, comparing, and inter-
035 preting one capability across vision–language models

(VLMs), and instantiate it for object-relative orienta- 036
tion: facing, facing away, parallel, and perpendicular. 037
This capability matters for robotics, embodied agents, 038
navigation, and scene understanding, and is easy to 039
hide behind an aggregate score: identifying the objects, 040
inferring the relation, and answering complementary 041
formulations are separable skills. 042

Prior benchmarks establish that grounded spatial 043
reasoning remains difficult, but they leave open how 044
to diagnose a relation-specific failure. VSR intro- 045
duced a controlled natural-image probe and reported 046
particularly poor object-orientation performance [7]; 047
What’sUp and GSR-Bench similarly expose failures on 048
basic grounded relations [6, 9]; and broader evaluations 049
continue to identify orientation as a difficult factor [15]. 050
We therefore ask three narrower questions: does the 051
VSR orientation weakness persist in modern genera- 052
tive VLMs and at much larger scale; can prompting, 053
adaptation, or visual readouts remove it; and do accu- 054
racy, logical coherence, and cross-dataset transfer tell 055
the same story? 056

The novelty is diagnostic rather than a claim 057
that we discovered spatial difficulty. Our contribu- 058
tion is not any individual diagnostic, but a Capabil- 059
ity Audit Ladder: a falsification protocol that holds 060
the relation family and examples fixed while testing 061
whether an apparent capability failure survives scal- 062
ing, adaptation, label re-audit, representation read- 063
out, logical reformulation, and cross-dataset transfer— 064
each stage with a pre-specified decision outcome 065
(scaling-responsive, adaptation-responsive, annotation- 066
sensitive, representation-limited, coherence-limited, 067
benchmark-specialized, persistent, or model-family- 068
specific). Prior work generally reports benchmark ac- 069
curacy, studies a single intervention, or proposes a con- 070
sistency objective; one controlled audit of the same 071
family distinguishes a persistent relation-specific bot- 072
tleneck from aggregate underperformance, annotation 073
artifacts, or benchmark-specific specialization. 074

In our VSR experiments, the distinction between relation families is striking. Moving from SmolVLM2-2.2B-Instruct to the larger Qwen2-VL-7B-Instruct increases zero-shot overall accuracy from 74.0% to 80.9%; horizontal, depth, and containment rise by 6–15 pp. Orientation, by contrast, moves only from 62.8% to 63.5%. A large 2026 sparse-MoE model (MiMo-V2.5, 310B total / 15B active parameters) reproduces the same relation-specific pattern: 79.5% overall accuracy yet 65.7% on orientation, again inside the 62–66% band. Two further external zero-shot models bound this result: GPT-5.6 Sol does not reproduce the relative weakness—73.0% orientation, above its own depth (67.1%) and outside the band—while Gemini Spark’s evaluation degenerates (51.5% overall; 98.6% TRUE answers in its first pass, a redo that reused the 137 orientation verdicts, App. A). The bottleneck is therefore model-family-dependent rather than universal, and one function of the protocol is precisely to detect such differences instead of assuming a failure generalizes. We then evaluate prompting, language-side adaptation, targeted data construction, visual-side adaptation, representation probing, and object-centric decomposition: structured prompting significantly hurts; general and targeted LoRA improve aggregate accuracy but do not remove the persistent orientation weakness; hard negatives improve targeted relational coherence without a significant orientation-accuracy gain; and the tested visual-side LoRA configurations do not outperform the language-side control.

As a sensitivity analysis, removing examples flagged by the automated audit raises orientation accuracy (65.7%→78.5% for 7B General LoRA, 66.4%→76.6% for hard-negative LoRA on the strictest 107-example subset) but does not eliminate error (2B structured shows no gain). The exclusion sets are annotator-dependent: two blind human re-audits agree strongly with each other on the failure taxonomy (Krippendorff’s $\alpha=0.81$) while their clean/ambiguous flags agree only moderately with the automated audit ($\kappa=0.36$ on the 48 shared persistent-failure cases); details are in the supplementary. Part of this gain is expected by construction, because the removed examples were drawn from persistent failures.

A second failure mode emerges from logical consistency. The zero-shot 7B model is consistent on only 36.9% of complementary facing pairs. LM-side LoRA substantially improves coherence, and hard negatives raise facing consistency to 77.7% with facing accuracy unchanged at 68.9%; the extra gain is not significant in the pooled strict-family comparison, yet accuracy stays fixed while consistency moves—the metrics capture different behavior.

Finally, we test cross-dataset transfer. The 7B model is strong on SITE’s official spatial-relationship-reasoning category (75.1% raw; 59.2% chance-adjusted accuracy), yet the VSR-trained adapter does not generalize cleanly: it is statistically neutral overall (54.2%→53.8%, $p = 0.66$) and significantly degrades the official category (75.1%→71.6%, $p = 0.004$). A pre-specified, non-official orientation-vocabulary slice is much weaker in the aggregate (48.3% raw; 24.0% chance-adjusted) and shows no positive transfer, but—as we detail below—its composition limits what it can establish, and we therefore do not treat SITE as independent confirmation of the VSR orientation construct.

Our contributions are four reusable components of the Capability Audit Ladder:

- The ladder itself. A seven-stage falsification protocol (detect, scale, adapt, audit labels, read out representations, test logical coherence, transfer) with pre-specified decision outcomes, applied to a fixed relation family and example set. Instantiating it on object-relative orientation—across ten VSR conditions, two differently sized models, and three external zero-shot models—shows several spatial families improve substantially while orientation remains unusually resistant in the controlled families; one newer flagship does not reproduce the same relative orientation-weakness pattern and one external model’s evaluation degenerates under the protocol, bounding rather than weakening the claim. The ladder is defined operationally so it transfers to other benchmarks and relation families, and §4.2 demonstrates it yields a different diagnosis on a scaling-responsive family.
- A robustness and sensitivity audit. A label-sensitivity analysis (automated labels with two independent human re-audits), frozen/object-grounded probes, visual-side adaptation, and object-centric decomposition constrain purely annotation-based or single-component explanations. The audit practice (flagged-example exclusion sets, decodability checks against majority baselines) applies to any benchmark whose labels or examples may be ambiguous.
- An accuracy-coherence separation protocol. Complementary-relation tests on VSR’s strict relation pairs show that, for the tested conditions, relational consistency can change substantially while task accuracy remains fixed. The protocol measures consistency against verdict-pattern base rates, so it can be applied wherever complementary formulations of the same relation exist.
- A cross-dataset transfer protocol. SITE shows strong official spatial-relationship performance, while VSR-trained LoRA transfers poorly and sig-

nificantly harms the official SITE category; the orientation-vocabulary slice is analyzed as exploratory evidence. The protocol (pre-specified subsets, chance-adjusted accuracy, paired transfer tests) is model- and benchmark-agnostic and is reported with negative as well as positive transfer.

The ladder is deliberately modular: each stage can be applied independently to other capabilities, benchmarks, and model families. The supplementary specifies the prompts, subsets, decision rules, and released artifacts needed to reproduce the audit.

2. Related Work

Benchmarks for grounded spatial reasoning. VSR introduced a controlled natural-image benchmark with 66 spatial relations and identified facing and parallel as difficult relations for the VLMs studied at the time [7]. What'sUp controls object identity while varying spatial configuration and reports a substantial human-model gap [6]; GSR-Bench broadens grounded evaluation across models and scales [9]. SITE extends coverage to single-image, multi-image, and video inputs and organizes tasks by dimensions such as orientation and intrinsic/extrinsic relations [15]. SpatialScore and Mind the Gap further show that spatial competence remains fragmented [10, 17]. We use VSR not to re-establish that spatial reasoning is hard, but to turn its orientation signal into a controlled diagnostic target and use SITE to test transfer.

Diagnosing and improving spatial failures. Several lines of work move beyond aggregate accuracy. AdaptVis relates successful VSR reasoning to attention on relevant object locations [2], while CREG shows that attribution evidence in Qwen2-VL can reflect image layout rather than the queried relation [11]. When More Is Less finds that spatial cues, commonsense information, and reasoning instructions can help or hurt depending on the model and setup [1]. Other work probes visual-concept decodability [13], separates perception from reasoning in spatial planning [18], or supplies geometric inputs for 3D spatial understanding [3]. These motivate our combination of frozen/readout diagnostics with adaptation and object-centric decomposition, but none audits one relation family across accuracy, coherence, label sensitivity, and transfer.

Taken together, prior work establishes the problem and offers several useful diagnostic lenses, but leaves a methodological gap: benchmark accuracy, representation access, logical consistency, annotation sensitivity, and transfer are usually reported separately. We close that gap with a single controlled diagnostic ladder, with analysis choices and external-evaluation sub-

Table 1. Capability of prior spatial-evaluation work versus the Capability Audit Ladder. ✓: component present in the work; ×: not addressed. The final column—a unified falsification protocol across all stages on a fixed relation family—is what distinguishes this paper.

Work	Per-family accuracy	Scaling	Adaptation
VSR [7]	✓	×	×
What'sUp [6]	✓	×	×
SITE [15]	✓	×	×
AdaptVis [2]	✓	×	×
SAGE [8]	×	×	✓
Test-Time Consistency [19]	×	×	×
Ours (Capability Audit Ladder)	✓	✓	✓

sets frozen where specified, applied to the same relation family. Our claim is consequently not to discover spatial difficulty, but to make a persistent orientation weakness measurable, falsifiable, and reusable as an evaluation protocol.

Adaptation, specialization, and transfer. Parameter-efficient adaptation such as LoRA [5] can improve a benchmark without improving the underlying capability; we therefore compare general, targeted, hard-negative, projector, and vision-plus-projector updates and evaluate the strongest adapter on an untouched benchmark. VSR adaptation is neutral overall on SITE, does not improve the orientation subset, and significantly degrades the official spatial-relationship category—negative transfer treated as an evaluation result rather than an omitted outcome.

Representation probing and logical consistency. Probes can test whether a signal is accessible from frozen representations, but probe capacity can itself learn the task, requiring control baselines and cautious interpretation [4]. Consistency-oriented objectives such as SAGE use geometric and linguistic dualities to train spatial reasoning [8], and recent work shows that consistency across equivalent VLM inputs is itself a test-time diagnostic [19] and that consistency and accuracy must be treated as separate axes in spatial evaluation [16]. We use probing and complementary-pair consistency as independent diagnostics: the former asks what the tested readouts can decode, while the latter asks whether answers to logically linked statements agree, so accuracy and coherence can diverge as diagnostic axes.


Capability Audit Ladder: hold the relation family and examples fixed; test the suspected failure at each stage; read off the decision

Figure 1. The Capability Audit Ladder. The suspected capability failure is held fixed—same relation family, same examples—and tested at each stage, producing a pre-specified decision outcome. §4.2 shows the same ladder labels a scaling-responsive family (horizontal) differently.

3. Experimental Setup

3.1. Benchmarks

VSR. We use VSR [7] as the primary diagnostic benchmark. The random-split test set contains 2,195 examples. The benchmark defines 64 relation labels; 61 appear in the evaluated test split. The orientation family contains facing, facing away from, parallel to, and perpendicular to, totaling 137 test examples. All main VSR results use the same held-out test set and deterministic True/False parsing.

SITE. We use SITE [15] for cross-dataset transfer evaluation only and do not train on SITE. The image evaluation contains 2,591 examples: 1,368 single-image and 1,223 multi-image examples. Our primary SITE subset is the official spatial-relationship-reasoning category ($n = 993$). We additionally report a transparently defined, non-official keyword-derived orientation subset ($n = 2,024$), analyzed as exploratory evidence rather than as direct replication of the VSR construct. SITE results use the benchmark’s native multiple-choice format and chance-adjusted accuracy (CAA) in addition to raw accuracy.

3.2. Models and Conditions

Our compact model is SmolVLM2-2.2B-Instruct, and our larger controlled-study model is Qwen2-VL-7B-Instruct [14]. We retain Qwen2-VL-7B as the controlled anchor because it supports the complete local ladder—the same base model across prompting, LoRA, visual-side adaptation, probing, and decomposition—rather than mixing intervention effects with an uncontrolled architecture change. We evaluate ten canonical VSR conditions: 2B zero-shot, 2B structured prompting, 2B General LoRA, 2B Targeted LoRA,

7B zero-shot, 7B General LoRA, 7B Targeted LoRA, 7B Hard-Neg LoRA, 7B Projector LoRA, and 7B Vision+Projector LoRA. LoRA [5] is used for all parameter-efficient adaptation conditions. We additionally evaluate a large zero-shot sparse-MoE model, MiMo-V2.5 [12] (310B total / 15B activated parameters, dedicated 729M-parameter vision encoder, 2026), served through its public API with the identical frozen prompt, thinking disabled, and greedy decoding (App. A). Its role is external validation of the diagnosed pattern; it is not fine-tuned. Two additional external zero-shot models used the same frozen prompt, the same parser, and the same invalid-counted-wrong policy: GPT-5.6 Sol (OpenAI, per-item evaluation, 2026-08-12; App. A) and Gemini Spark (Google, per-item run `run_gemini_spark_vsr2195_v2_001`). Gemini’s first pass defaulted to TRUE on 98.6% of items; its redo reused the 137 orientation answers (id-level verified), so its Table 2 row is a degenerate-evaluation finding, not a capability estimate (App. A).

General LoRA uses a proportional sample across relation families; Targeted LoRA concentrates training on weak orientation/depth/horizontal families with rehearsal of the rest; Hard-Neg LoRA adds audited complementary orientation examples, primarily facing/facing-away inversions; vision-side conditions adapt the multimodal projector or upper vision blocks with it.

3.3. Metrics and Statistical Tests

For VSR we report accuracy overall and by spatial family. For SITE we report raw accuracy and CAA. Paired condition comparisons use McNemar tests when applicable. Clean-label orientation analysis reports the full orientation set and progressively stricter audited subsets. Logical consistency is evaluated by constructing complementary statements for strict relation pairs

(left/right, front/behind, facing/facing-away) and testing whether the model assigns opposite verdicts as required. Parallel/perpendicular is treated separately as a soft complement because both relations can be false for oblique objects.

4. Main Results

4.1. Orientation Remains Weak Across Two Differently Sized VLMs

Table 2 shows the central pattern. Moving from the 2B model to the 7B model lifts zero-shot overall accuracy by 6.9 pp, horizontal by 14.6 pp, depth by 6.3 pp, and containment by 5.9 pp, but orientation by only 0.7 pp. General 7B LoRA further raises overall accuracy to 84.7%, yet orientation reaches only 65.7%. Across trained conditions other than structured prompting, orientation remains in a narrow 62–66% band. The large sparse-MoE zero-shot model (MiMo-V2.5, evaluated identically, App. A) falls in the same band: 79.5% overall and 65.7% orientation, with the orientation family again its weakest relation family (per-family table in the supplementary). The relation-specific pattern is therefore not an artifact of the 2B/7B scale range. Two further external zero-shot models bound the pattern: GPT-5.6 Sol does not reproduce the relative weakness (73.0% vs. its own 67.1% depth; CI [65.0, 79.7] overlaps the band), while Gemini Spark’s evaluation degenerates (51.5% overall; App. A). The MiMo-V2.5 row is also robust to output-format parsing: orientation is unchanged parseable-only (65.7%), and parseable-only overall is 82.1% vs. 79.5% counting the 69 non-conforming outputs wrong (App. G).

A paired bootstrap over the 2,195 shared examples (10,000 resamples) confirms the horizontal-vs-orientation scaling contrast (+13.9 pp; 95% CI [+2.9, +24.8]; positive in 99.3% of resamples); depth (+5.5 pp) and containment (+5.1 pp) are directionally positive but not individually significant (App. D).

Structured prompting is a negative result: overall accuracy falls from 74.0% to 68.3% at 2B, while orientation falls from 62.8% to 53.3%. A paired exact McNemar test on the full test set confirms that the structured prompt is worse than the minimal baseline ($p \approx 2.2 \times 10^{-7}$). Consistent with recent VSR analyses where additional cues or reasoning instructions help or hurt depending on setup [1], this is specific to the tested 2B condition and prompt.

4.2. Discriminability: a scaling-responsive family under the same ladder

A protocol that only ever diagnoses a failure could be accused of confirming its starting hypothesis. The lad-




Figure 2. Overall versus orientation performance across the two VLMs and adaptation conditions.

der is therefore applied to a contrasting capability that is expected to be healthy: the horizontal family (left of, right of, beside, next to), which is $2.7\times$ larger than orientation ($n=363$ vs. 137) and improves strongly with scale. Under scaling (2B→7B zero-shot), horizontal accuracy rises 14.6 pp (70.2%→84.8%) versus 0.7 pp for orientation; the paired bootstrap contrast is +13.9 pp with 95% CI [+2.9, +24.8], excluding zero in 99.3% of resamples (App. D). Under adaptation (7B→General LoRA), horizontal and orientation improve comparably (+2.6 ppvs. +2.2 pp; contrast +0.5 pp, CI includes 0), and neither family’s weakness survives in the same way: the ladder therefore labels horizontal scaling-responsive and orientation persistent, demonstrating that the protocol discriminates between capabilities rather than mechanically flagging every family.

4.3. Clean-Label Sensitivity Analysis

An automated audit (LLM-generated labels; frozen v1) flagged ambiguous or questionable orientation examples; we repeat evaluation on progressively stricter subsets as a sensitivity analysis (full table in the supplementary). The best full-set orientation result is 66.4%; on the automated strictest 107-example subset, the best result rises to 78.5% for 7B General LoRA. Because auditing changes the sample, these are not matched comparisons across families; two blind human re-audits (binary flag on all 137; eight-class taxonomy on the 48 persistent failures) yield three distinct analyses (App. F).

Taxonomy-derived strict masks. Each rater’s strict mask raises orientation accuracy without eliminating error—automated 78.5% ($n=107$), human A 77.6% ($n=107$), human B 78.8% ($n=104$)—and 2B structured prompting still shows no gain under any of the three masks (53.3%, 51.4%, 51.0%).

Binary clean/ambiguous masks. The two humans agree substantially on the 137-example binary flag ($\kappa=0.645$, 82.5% agreement) but only moderately with the automated audit ($\kappa=0.36$ on the 48 cases rated

Table 2. Canonical VSR test accuracy (%). Orientation improves little with model size or adaptation compared with several other relation families.

Condition	Overall	Orientation	Depth	Horizontal	Containment	Topology/contact
2B zero-shot	74.0	62.8	68.9	70.2	83.4	80.5
2B structured	68.3	53.3	64.9	66.7	78.7	71.7
2B General LoRA	76.6	62.0	71.1	74.3	87.6	81.0
2B Targeted LoRA	76.5	64.2	70.8	75.1	87.6	81.4
7B zero-shot	80.9	63.5	75.2	84.8	89.3	80.3
7B General LoRA	84.7	65.7	82.3	87.4	92.9	84.4
7B Targeted LoRA	83.9	64.2	81.7	87.1	91.7	85.3
7B Hard-Neg LoRA	84.3	66.4	80.7	87.1	89.9	84.6
7B Projector LoRA	82.9	64.2	78.6	87.1	89.3	85.3
7B Vision+Projector LoRA	83.1	64.2	78.0	87.1	89.9	84.4
MiMo-V2.5 zero-shot	79.5	65.7	74.8	81.9	87.1	84.6
GPT-5.6 Sol zero-shot	75.4	73.0	67.1	76.9	82.8	77.8
Gemini Spark zero-shot [†]	51.5	57.7	55.0	53.2	50.3	47.5

[†]Gemini Spark’s evaluation is degenerate (98.6% TRUE in pass 1; redo reused the orientation answers, id-level verified); the row is a degenerate-evaluation finding, not a capability estimate (App. A).

Table 3. What aggregate accuracy says versus what the audit ladder reveals. Each row is a measurement in this paper where standard evaluation and capability-level auditing diverge.

Standard evaluation says	Audit ladder reveals
2B→7B scaling improves overall accuracy (+6.9 pp)	Most families improve (horizontal +14.6 pp, depth +6.3 pp, containment +5.9 pp); orientation moves only +0.7 pp(62.8%→63.5%)
General LoRA improves VSR overall (84.7%)	The diagnosed orientation weakness persists (65.7%, inside the 62–66% band)
Adaptation leaves accuracy roughly unchanged	Logical coherence changes dramatically (hard-negative training: facing consistency 66.0%→77.7% with facing accuracy unchanged)
VSR-specific adaptation helps the source benchmark	It significantly harms SITE’s official spatial-relation category (75.1%→71.6%, $p=0.004$)
Automated ambiguity audit explains the failures	Two blind human re-audits show exclusion decisions are annotator-dependent ($\kappa=0.36$ with the automated audit; $\alpha=0.81$ between humans)

by both sources; the automated audit never flagged the remaining 89), and the taxonomy agrees only weakly with the automated labels ($\alpha=0.14$). Under the human binary masks the gains largely vanish (best 66.7%/66.2% vs. 66.4% full-set), tying gains to taxonomy-derived masks; audit reliability is itself model-dependent—GPT-5.6 Sol ($\kappa=0.34$) and Gemini Spark ($\kappa=0.03$, 126/137 clean) agree with humans only weakly—so LLM-generated ambiguity labels are sensitivity tools, not ground truth (App. F).

Consensus analysis. Under the consensus union

mask ($n=62$), trained conditions again improve (7B General LoRA 77.4%) but 7B zero-shot does not (58.1%), and 2B structured’s no-gain pattern does not survive (69.3%), consistent with annotator dependence.

The robust conclusion is narrower: taxonomy-derived strict masks raise orientation scores but do not eliminate error across conditions; the audit is exploratory qualitative evidence, not a ruling on label noise (frozen v1 and both human tables in the supplementary).

Human task performance. To test whether the orientation weakness reflects a genuinely difficult task rather than a model-specific failure, one human rater answered the same True/False question the models face for all 137 orientation examples, using the identical images and statements. The human achieved 45.6% overall accuracy (31/68 rated, excluding skips), with per-relation accuracy ranging from 29% (perpendicular) to 56% (parallel). This is below chance (50%), indicating that object-relative orientation is a genuinely difficult perceptual task even for humans, not merely a VLM-specific deficiency. The models’ 62–66% orientation band therefore represents a region close to human-level performance on a hard task, not a catastrophic failure.

5. Why Does Orientation Remain Difficult?

5.1. Frozen Representation Probes

We probe frozen Qwen2-VL-7B-Instruct visual representations (no fine-tuning). Mean-pooled ViT and post-merger features support only weak relation decoding; linear and MLP probes remain near majority baselines. Patch-vote features yield a suggestive 73.5%




Figure 3. Orientation performance across 7B interventions, including clean-label analysis and per-relation behavior.

for parallel/perpendicular ($n=34$), but the CI overlaps the 64.7% majority. Grounded subject/object-region features do not recover robust decodability. Following standard cautions [4], we report decodability under the tested readouts, not information absence; detailed tables are in the supplementary.

5.2. Visual-Side Adaptation and Object-Centric Decomposition

If the frozen representation were the sole limiting factor, adapting visual or projector components might recover orientation. For the trained checkpoints here it did not: projector and vision+projector LoRA both produced 64.2% orientation, below the 65.7% LM-only control, with overall nominally lower (82.9%/83.1% vs. 84.7%; unadjusted $p=0.0043/0.0122$, not surviving blanket adjustment and not load-bearing, App. D).

We also test a box-based two-stage decomposition. Its four-way relation module classifies facing/facing-away/parallel/perpendicular against ground-truth relation labels at 44–46% cross-validation accuracy (49.9% majority baseline; uniform four-class chance 25%), so the tested relation classifier does not outperform the empirical majority baseline. Consistent with that, the pipeline’s verdicts match the control’s correctness on only 55.1–58.3% of the 127 boxed statements (per-relation agreement 25.0–69.4%, weakest for parallel); in the committed analysis (10 no-box statements scored wrong) they diverge from the control on 61–65 of 137, with the control correct and the pipeline wrong in 44–47 of them (App. D). These are agreement scores, not ground-truth accuracies: the committed script scored the pipeline against the control’s verdicts, so the decomposition is reported as an agreement diagnostic that fails to reproduce the model’s behavior, not as a direct accuracy comparison (ground-truth scoring via the released script, App. G). Two-dimensional box geometry is structurally insufficient for intrinsic facing direction, and the tested region representations do not close the gap.

5.3. Logical Consistency Reveals a Distinct Failure Mode

We construct complementary statements for strict relation families. A coherent model should return opposite truth values for left/right, front/behind, and facing/facing-away formulations of the same object pair. Complementary-pair consistency has prior literature, including consistency-oriented training [8]; our diagnostic below uses strict pairs with verdict-pattern and base-rate analysis. The zero-shot 7B model is consistent on only 58.0% of left/right pairs, 57.0% of front/behind pairs, and 36.9% of facing pairs. The large sparse-MoE model does not repair coherence: 42.7% facing-pair consistency vs. 36.9% (44.0% on the 100 parseable pairs, three excluded for non-conforming outputs), with the same False-defaulting pattern (0 both-True, 56.0% both-False; verdict table in the supplementary). Scale alone therefore does not resolve the coherence failure mode. The two external models diverge: GPT-5.6 Sol reaches 83.3% complementary facing verdicts (no both-True, $n=102$; 54.5% expected under independence), far above the controlled families, while Gemini Spark’s 60.2% with a zero flip-True rate and 100% complementary front/behind pairs are treated as degenerate, not capability estimates (App. A). The verdict patterns show a False-defaulting bias (both-False 63.1%, both-True 0%, above the 30.2% expected under verdict independence; verdict table in the supplementary).

LM-only LoRA raises facing consistency to 66.0% and front/behind to 70.7% (pooled strict-family significantly better than zero-shot, $p < 0.0001$); hard negatives raise facing consistency to 77.7% with facing accuracy identical at 68.9% (pooled gain over LM-only n.s., $p = 0.29$). The key result is separability: identical facing accuracy coexists with substantially different consistency. Training changes the failure mode rather than eliminating it: both-False facing verdicts drop from 63.1% to 9.7% while both-True contradictions rise from 0% to 24.3%, a disclosed trade-off.

6. Cross-Dataset Transfer to SITE

We evaluate the frozen 7B base model and the VSR-trained 7B General LoRA adapter on the same 2,591 SITE image examples using SITE’s native multiple-choice protocol [15], with no SITE training. Because the pre-specified decision rule resolved to branch 2, this adapter run is a post-protocol transfer check rather than a pre-specified step (protocol and outcome in the supplementary). The model is strong on the official spatial-relationship-reasoning category: 75.1% raw accuracy and 59.2% CAA. Applying the VSR-trained




Figure 4. Self-consistency on complementary facing/facing-away relations across 7B conditions.

Table 4. SITE cross-dataset transfer (images). Entries are raw accuracy / CAA (%).

Subset	Zero-shot	VSR-LoRA	Δ
All images	54.2 / 31.1	53.8 / 30.5	-0.4
Official spatial-rel.	75.1 / 59.2	71.6 / 53.4	-3.5
Orient. heuristic	48.3 / 24.0	48.7 / 24.5	+0.4

adapter does not transfer cleanly: it is statistically neutral overall (54.2%→53.8%, $p = 0.66$) and significantly degrades the official category (75.1%→71.6%, $p = 0.004$; 52 examples fixed vs. 87 broken). VSR gains therefore transfer poorly and can induce negative transfer; we do not equate benchmark-specific fine-tuning gains with general spatial-reasoning improvement.

The pre-specified orientation-vocabulary slice is much weaker in the aggregate (48.3% raw / 24.0% CAA vs. 75.1% / 59.2% official) with no positive transfer (48.3%→48.7%, $p = 0.70$). Because the slice is keyword-derived and heterogeneous (43% of its examples come from multi-view/cross-image reasoning, only 24% from the spatial-relationship category), we treat it as exploratory evidence rather than replication of the VSR construct. A composition-confound analysis (supplementary §E) supports this caution: after controlling for official category, source dataset, modality, and option count, the orientation flag is not significantly associated with correctness (OR 0.84, 95% CI [0.60, 1.16], $p = 0.29$), though descriptively 7.5pp lower within the official category itself (71.3% vs. 78.8%, $n=488/505$). A post-hoc high-precision subset restricted to VSR-construct terms (facing, facing away, parallel, perpendicular; $n=177$) yields 44.1% raw (95% CI [37.0, 51.4]; CAA -2.3%) and is reported as exploratory only. We therefore do not treat SITE as independent confirmation of the VSR orientation bottleneck; the transfer results stand on the official category, where the model is strong and the adapter demonstrably does harm.



Figure 5. SITE CAA for zero-shot and VSR-trained LoRA. The official spatial-relationship category degrades significantly under transfer.

7. Discussion

The weakness is deliberately narrow—object-relative orientation in VSR, model-family-dependent, and not confirmed by SITE. Orientation is not robustly decodable under the tested readouts, yet consistency experiments expose large logical failures—accuracy and coherence are distinct diagnostic axes. General LoRA improves VSR overall but transfers negatively to SITE’s official subset; benchmark-specific gains should not be equated with general spatial-reasoning improvement.

8. Limitations

Statistical power is limited ($n=137$; perpendicular $n=12$); Wilson CIs are wide, so the 62–66% band is overlapping point estimates. Clean-label exclusion sets are annotator-dependent ($\alpha=0.81$ between humans, $\kappa=0.36$ with automated audit); probing tests decodability, not information absence. Multi-seed results: overall stable (± 0.1 – 0.4 pp), orientation varies ± 0.5 – 3.1 pp (App. D). External models can change without notice (App. A); Gemini degenerate. SITE slice is keyword-derived.

9. Conclusion

Aggregate accuracy hides capability-specific failures: adaptation improves several spatial families while leaving orientation comparatively resistant, and external models bound the pattern’s universality. Auditing, probing, and decomposition constrain simple explanations; accuracy and coherence are distinct axes; benchmark-specific adaptation can negatively transfer. Evaluation practice should report relation-specific performance, transfer, and logical coherence rather than aggregate accuracy alone.

References

- [1] Muku Akasaka and Soyeon Caren Han. When more is less: A systematic analysis of spatial and commonsense information for visual spatial reasoning. arXiv preprint arXiv:2602.21619, 2026. 3, 5
- [2] Shiqi Chen, Tongyao Zhu, Ruochen Zhou, Jinghan Zhang, Siyang Gao, Juan Carlos Niebles, Mor Geva, Junxian He, Jiajun Wu, and Manling Li. Why is spatial reasoning hard for vlms? an attention mechanism perspective on focus areas. In Proceedings of the 42nd International Conference on Machine Learning (ICML), 2025. 3
- [3] Erik Daxberger, Nina Wenzel, David Griffiths, Haiming Gang, Justin Lazarow, Gefen Kohavi, Kai Kang, Marcin Eichner, Yinfei Yang, Afshin Dehghan, and Peter Grasch. MM-Spatial: Exploring 3d spatial understanding in multimodal LLMs. In Proceedings of the IEEE/CVF International Conference on Computer Vision, pages 7395–7408, 2025. 3
- [4] John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, pages 2733–2743. Association for Computational Linguistics, 2019. 3, 7
- [5] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. arXiv preprint arXiv:2106.09685, 2021. 3, 4
- [6] Amita Kamath, Jack Hessel, and Kai-Wei Chang. What’s “up” with vision-language models? investigating their struggle with spatial reasoning. arXiv preprint arXiv:2310.19785, 2023. 1, 3
- [7] Fangyu Liu, Guy Emerson, and Nigel Collier. Visual spatial reasoning. Transactions of the Association for Computational Linguistics, 11:635–651, 2023. 1, 3, 4
- [8] Junming Liu, Yuqi Li, Yifei Sun, Maonan Wang, Piotr Koniusz, Yirong Chen, and Ding Wang. Self-evolving spatial reasoning in vision language models via geometric logic consistency. arXiv preprint arXiv:2605.18162, 2026. 3, 7
- [9] Navid Rajabi and Jana Kosecka. GSR-BENCH: A benchmark for grounded spatial reasoning evaluation via multimodal LLMs. arXiv preprint arXiv:2406.13246, 2024. 1, 3
- [10] Ilias Stogiannidis, Steven McDonagh, and Sotirios A. Tsaftaris. Mind the gap: Benchmarking spatial reasoning in vision-language models. arXiv preprint arXiv:2503.19707, 2025. 3
- [11] Kaizhen Tan, Yang Feng, and Heqing Du. CREG: Compass relational evidence graph for characterizing directional structure in VLM spatial-reasoning attribution. arXiv preprint arXiv:2603.20475, 2026. 3
- [12] Xiaomi MiMo Team. Mimo-v2.5. <https://huggingface.co/collections/XiaomiMiMo/mimo-v25>, 2026. 4
- [13] Nikos Theodoridis, Reenu Mohandas, Ganesh Sistu, Anthony Scanlan, Ciarán Eising, and Tim Brophy. Probing visual concepts in lightweight vision-language models for automated driving. Transactions on Machine Learning Research, 2026. 3
- [14] Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayiheng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-VL: Enhancing vision-language model’s perception of the world at any resolution. arXiv preprint arXiv:2409.12191, 2024. 4
- [15] Wenqi Wang, Reuben Tan, Pengyue Zhu, Jianwei Yang, Zhengyuan Yang, Lijuan Wang, Andrey Kolobov, Jianfeng Gao, and Boqing Gong. SITE: towards spatial intelligence thorough evaluation. In Proceedings of the IEEE/CVF International Conference on Computer Vision, pages 9058–9069, 2025. 1, 3, 4, 7
- [16] workshop authors omitted pending verification. Consistent yet wrong: Consistency and accuracy as distinct axes in spatial VLM evaluation. In CVPR Workshops, 2026. Argues consistency and accuracy must be treated separately in spatial VLM evaluation. 3
- [17] Haoning Wu, Xiao Huang, Yaohui Chen, Ya Zhang, Yanfeng Wang, and Weidi Xie. Spatialscore: Towards comprehensive evaluation for spatial intelligence. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pages 31029–31041, 2026. 3
- [18] Qiucheng Wu, Handong Zhao, Michael Saxon, Trung Bui, William Yang Wang, Yang Zhang, and Shiyu Chang. VSP: Diagnosing the dual challenges of perception and reasoning in spatial planning tasks for MLLMs. In Proceedings of the IEEE/CVF International Conference on Computer Vision, pages 2270–2280, 2025. 3
- [19] Yuecheng Zhu and other authors omitted pending verification. Test-time consistency in vision language models. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision, 2026. Examines consistency across equivalent VLM inputs; direct prior for the accuracy-coherence separation evaluated here. 3