

Beyond Spatial Accuracy: Diagnosing Persistent Orientation Failures in Vision–Language Models

Anonymous WACV [Datasets Track](#) submission

Paper ID *****

Beyond Spatial Accuracy: Diagnosing Persistent Orientation Failures in Vision–Language Models

Supplementary Material

001	1. Exact Prompts, Decoding, and Parsing	045
002	VSR plain prompt (identical for 2B and 7B, all condi-	046
003	tions; '{statement}' is the VSR caption):	047
004	Look at the image carefully.	048
005		049
006	Statement: "{statement}"	050
007		051
008	Is this statement true or false?	052
009		053
010	Answer with exactly one word:	054
011	True or False.	055
012	2B structured-decomposition prompt (2B only; sin-	056
013	gle run):	057
014	Look carefully at the image.	
015		
016	Statement:	
017	"{statement}"	
018		
019	Determine whether the statement is true	
020	by doing the following:	
021	1. Identify the subject object.	
022	2. Identify the reference object.	
023	3. Determine the position and	
024	orientation of the subject relative	
025	to the reference object.	
026	4. Pay careful attention to left/right,	
027	front/back, viewpoint, and	
028	orientation.	
029	5. Decide whether the stated spatial	
030	relationship matches the image.	
031		
032	Answer with exactly one word:	
033	True or False.	
034	SITE multiple-choice prompt (replicated from the	
035	official lmms-eval task sitebench/utlis.py):	
036	# image examples (one <image> token per visual):	
037	<image>...	
038	Question: {question}	
039	Options:	
040	A: {opt0}	
041	B: {opt1}	
042	...	
043	Give me the answer letter directly.	
044	The best answer is:	
	# video examples (pre-prompt):	
	Select the best answer to the following	
	multiple-choice question based on the	
	video. Respond with only the letter	
	of the correct option.	
	Question: {question}	
	Options:	
	A: ...	
	B: ...	
	...	
	Give me the answer letter directly.	
	The best answer is:	
	Decoding. All generations are greedy:	
	do_sample=False, temperature 0, top-p 1.0.	
	max_new_tokens: 5 for the 2B wrapper, 128 for	
	7B VSR runs, 16 for SITE (change validated parse-	
	identical on 125/125 held-out outputs before adoption;	
	recorded in the SITE run metadata).	
	MiMo-V2.5 protocol (external-validation row).	
	MiMo-V2.5 (XiaomiMiMo/MiMo-V2.5, 310B total	
	/ 15B activated, dedicated 729M vision encoder) is	
	evaluated zero-shot through its public API with the	
	identical frozen VSR prompt above, thinking mode dis-	
	abled (chat_template_kwargs.enable_thinking=false	
	and thinking.type=disabled; verified content-only,	
	zero reasoning tokens, deterministic on re-runs),	
	temperature 0, top-p 1.0, max_tokens=64, one image	
	per request (canonical COCO URLs, with automatic	
	base64 retry fallback from the local cache), and 12	
	parallel workers (within the provider's rate limits).	
	Outputs use the same parser; 3.1% of responses	
	did not conform to the one-word format and are	
	counted as wrong (never guessed), and 294 responses	
	contained chain-of-thought text alongside a parseable	
	verdict; both are flagged in the committed predictions	
	(results/mimo/). The evaluation cost is reported in	
	results/mimo/usage_report.json (per-request token	
	accounting).	
	Parsing. VSR outputs are parsed with an anchored	
	case-insensitive regex (src/evaluation/parser.py):	
	strips prefixes ('the answer is', 'answer:', 'response:',	
	'output:') and trailing punctuation; accepts exact	
	'true/yes/correct' or 'false/no/incorrect' forms; sub-	
	string fallback only when exactly one of 'true'/'false'	
	appears. Invalid outputs are counted as wrong, never	
	guessed; invalid rate is 0 on all reported runs. SITE	

092 outputs use the official parse_multi_choice_response
093 (letter); unparseable responses are counted incorrect
094 and reported.

095 2. Training and Adapter Configurations

096 All LoRA conditions share: $r=8$, $\alpha=16$, dropout
097 0.05, AdamW lr 10^{-4} , weight decay 0.01, linear
098 warmup 10% of steps, 2 epochs, bf16, gradient
099 checkpointing, gradient clipping 1.0, vision tower
100 frozen. Adapter configs are committed at check-
101 points/*/final/adapter_config.json.

102 Manifest construction. General: 2,000 state-
103 ments sampled from the VSR train split. Tar-
104 getted: orientation-over-sampled subset of the same
105 split. Hard-negative: the orientation block of the
106 general manifest is replaced by audited-clean originals
107 plus paired hard negatives (facing $\leftrightarrow$ facing away from;
108 parallel $\leftrightarrow$ perpendicular) with flipped labels; every in-
109 cluded example was audited for label correctness (451
110 train orientation examples audited: 402 clean, 28 ex-
111 cluded, 21 original-only). 2B runs split 5% of the man-
112 ifest for validation (seed 42). All manifests derive ex-
113 clusively from the VSR train split; all evaluations use
114 test only.

115 3. Probe Methodology and Full Tables

116 Features. Qwen2-VL-7B-Instruct, vision tower frozen.
117 ViT patch embeddings (1280-d) and post-merger fea-
118 tures (3584-d), mean-pooled per image (ungrounded)
119 or per grounding box (grounded). Grounding boxes
120 from frozen Qwen2-VL grounding (94.6% of 647 ori-
121 entation images yield both boxes; failures are mostly
122 genuinely absent objects), rescaled to pixels; region fea-
123 tures mean-pooled over patch centers inside each box.

124 Tasks. T1 facing vs. facing away (test $n=103$, major-
125 ity 64.2%); T2 parallel vs. perpendicular ($n=34$, major-
126 ity 64.7%); T3 4-way ($n=137$, majority 49.6%). Train-
127 ing on audited-clean VSR train; 5-fold CV; no test ex-
128 amples in training.

129 Readouts. Linear logistic regression; MLP (256
130 hidden, ReLU, Adam, early stopping on CV); patch-
131 vote (per-patch linear probe, majority vote over
132 patches); grounded variants with feature sets visual
133 ([subj, ref, subj-ref, subj-ref]), geometry (12-d normal-
134 ized box features), and visual+geometry.

135 Ungrounded results (test acc %; balanced acc in
136 parentheses; Wilson 95% CI; majority baselines: T1
137 64.2%, T2 64.7%, T3 49.6%):

T1 facing vs. away ($n=103$)	Test	Bal. acc	95% CI	
ViT linear	65.0	60.4	[55, 74]	138
ViT MLP	55.3	52.0	[45, 65]	
Merger linear	68.9	64.5	[59, 77]	
Merger MLP	52.4	49.2	[42, 62]	
ViT patch-vote	62.1	—	[53, 71]	
Merger patch-vote	60.2	—	[51, 69]	

139

T2 parallel vs. perpendicular ($n=34$)	Test	Bal. acc	95% CI	
ViT linear	61.8	61.0	[44, 77]	140
ViT MLP	67.6	69.3	[50, 82]	
Merger linear	64.7	61.4	[47, 79]	
Merger MLP	64.7	63.3	[47, 79]	
ViT patch-vote	73.5	—	[57, 85]	
Merger patch-vote	67.6	—	[51, 81]	

T3 4-way ($n=137$)	Test	Bal. acc	95% CI	
ViT linear	52.6	40.5	[44, 61]	141
ViT MLP	52.6	42.9	[44, 61]	
Merger linear	53.3	44.0	[45, 61]	
Merger MLP	47.4	43.2	[39, 56]	
ViT patch-vote	51.8	—	[44, 60]	
Merger patch-vote	52.6	—	[44, 61]	

Grounded results (CV acc / test acc; majority T1 142
63.7%, T2 57.0%; visual = [subj, ref, subj-ref, subj-ref]; 143
geometry = 12-d box features; T3 4-way is at chance 144
for every readout, test acc 41.7–52.8% across all 12 cells 145
vs. 49.9% majority, committed JSON): 146

Readout	T1 visual	T1 geom.	T2 visual	T2 geom.	
ViT linear	59.8/61.6	61.1/60.6	51.2/53.6	44.2/50.0	148
ViT MLP	59.5/66.7	61.8/62.6	45.4/46.4	45.4/46.4	
Merger linear	57.2/65.7	61.1/60.6	51.2/53.6	44.2/50.0	
Merger MLP	56.3/52.5	61.1/71.7	49.0/53.6	50.0/42.9	

The single above-majority cell (geometry-only merger 149
MLP on T1, 71.7% test) has CV 61.1% below the 150
63.7% majority — not robust — and derives from 151
the grounding model’s own box statistics, not vision- 152
feature content. In short, object-box conditioning 153
with the model’s own grounding predictions and mean- 154
pooled region features does not recover robust orien- 155
tation decodability under the tested readouts. The 156
T2 patch-vote result (73.5%) is a suggestive point esti- 157
mate with a wide CI overlapping the majority baseline 158
($n=34$). 159

Clear-subset control (canonical frozen snapshot). 160
The 48-example automated (LLM-generated) audit 161
identified ambiguous or questionable orientation la- 162
bels; its reliability is assessed against an inde- 163
pendent human re-audit in §6. Re-evaluation on 164
progressively stricter subsets (values % from re- 165
sults/tables/orientation_clean_label_table.md, gen- 166
erated by scripts/clean_label_orientation.py): Be- 167
cause auditing changes the evaluated sample, strict- 168
subset values are not a matched comparison against 169

Condition	LoRA targets	Trainable	Batch	Manifest
2B General	LM q/k/v/o	~2.3M	microbatch 4×4	general (2,000)
2B Targeted	LM q/k/v/o	~2.3M	same	targeted (2,000)
7B LM-only	LM q/k/v/o	~0.3M	1	general (2,000)
7B HardNeg	LM q/k/v/o	~0.3M	1	hard-negative (2,000)
7B Projector	visual.merger.mlp.0/2	152K	1	general (2,000)
7B Vision+Proj	merger + blocks 24–31	1.46M	1	general (2,000)

Table 1. LoRA conditions and their target modules.

Table 2. Clean-label orientation robustness (canonical frozen snapshot; values % from re-
sults/tables/orientation_clean_label_table.md).

Condition	Full (137)	–Questionable (132)	Clear (124)	Strict (107)
2B zero-shot	62.8	65.2	66.9	75.7
2B structured	53.3	52.3	51.6	53.3
2B General LoRA	62.0	63.6	64.5	72.9
2B Targeted LoRA	64.2	65.9	68.5	76.6
7B zero-shot	63.5	65.2	68.5	72.0
7B General LoRA	65.7	67.4	70.2	78.5
7B Targeted LoRA	64.2	65.9	69.4	72.0
7B Hard-Neg LoRA	66.4	68.2	70.2	76.6
7B Projector LoRA	64.2	65.9	69.4	73.8
7B Vision+Projector LoRA	64.2	65.2	68.5	73.8
MiMo-V2.5 zero-shot	65.7	67.4	69.3	71.0

other relation families. The earlier probe-side clear/strict check (pre-freeze) additionally suggested that removing audit-flagged examples does not change the probe conclusion: the best ungrounded probe moved only 68.9% → 71.1% → 71.6% (full/clear/strict, $n = 103/97/88$) and the generative T1 readouts 64.1/66.0/68.2 (zero-shot) and 68.9/72.2/78.4 (General LoRA); committed in results/probe/clear_subset_results.json. The reliability of the automated audit underlying this table is assessed in §6 (human re-audit + inter-source agreement); the frozen table above is retained unchanged.

Human re-audit table (versioned, additive). The exclusion masks of Table 2 were originally produced by an automated (LLM-generated) audit whose image grounding is unverified. An independent human annotator re-rated all 137 orientation examples (binary clean/ambiguous flag) and the 48 persistent failures (eight-class taxonomy) from the images alone. Applying the same taxonomy-derived strict mask to the human labels (again $n=107$) reproduces the qualitative conclusions, with lower ceilings:

Under the human strict mask, orientation accuracy rises (best 77.6%, 2B Targeted LoRA), error is not eliminated (human ceiling 76.6% below the automated ceiling 78.5% for the 2B/7B conditions), and 2B structured still shows no gain (51.4%), so the paper’s sensitivity conclusions are not an artifact of the automated audit. The human binary mask alone (75 clean examples)

does not reproduce the gains (54.7–66.7%), which we read as further evidence that clean-label judgments are labeler-dependent at the per-item level; per-item agreement and a consensus analysis are in §6.

4. Complete Statistical Test Battery

All pairwise tests are exact McNemar tests (binomial test on discordant pairs, two-sided). Discordant counts are reported as (condition-loses / condition-gains) relative to the control named in each row. No test was run and withheld; the battery below is complete for this paper.

Principal vs. exploratory comparisons. The battery above is fully enumerated; no test was run and withheld. The principal (confirmatory) comparisons are: zero-shot vs. LM-only consistency (ID 19) and the SITE transfer tests (§E). All other comparisons — vision-side conditions and their per-relation cells, hard-negative training vs. control, the remaining consistency comparisons, and the two-stage agreement rows (IDs 15–18) — are exploratory diagnostics reported with exact unadjusted p -values. Under a Holm adjustment over the 18 accuracy-valid VSR tests ($\alpha=0.05$), only the zero-shot-vs-LM-only consistency test survives (Holm $p=0.0018$, using the conservative $p<0.0001$ bound); the vision-side overall effects (raw $p=0.0043$, $p=0.0122$; Holm 0.073, 0.195) and all other comparisons do not survive a blanket adjustment and are not load-bearing. For the 3 SITE transfer tests, Holm-

Table 3. Human re-audit clean-label orientation robustness (versioned, additive; values % from results/tables/orientation_clean_label_table_human.md). The frozen automated-audit table (Table 2) is not modified.

Condition	Full (137)	−Q	Clear	Strict (107)	Human binary (75)
2B zero-shot	62.8	67.5	67.5	76.6	65.3
2B structured	53.3	53.2	53.2	51.4	66.7
2B General LoRA	62.0	66.7	66.7	74.8	66.7
2B Targeted LoRA	64.2	69.0	69.0	77.6	65.3
7B zero-shot	63.5	65.9	65.9	74.8	54.7
7B General LoRA	65.7	68.2	68.2	72.9	66.7
7B Targeted LoRA	64.2	65.1	65.1	69.2	57.3
7B Hard-Neg LoRA	66.4	68.2	68.2	72.9	65.3
7B Projector LoRA	64.2	65.1	65.1	71.0	61.3
7B Vision+Projector LoRA	64.2	67.5	67.5	72.9	64.0
MiMo-V2.5 zero-shot	65.7	68.2	68.2	72.0	58.7

ID	Comparison	Subset	Discord.	<i>p</i>
1	HardNeg vs. Gen.	overall (2195)	52/60	0.508
2	HardNeg vs. Gen.	weak families	20/26	0.461
3	Projector vs. ctrl.	overall	—	0.0043
4	V+Proj vs. ctrl.	overall	—	0.0122
5	Projector vs. ctrl.	orientation	—	0.86
6	V+Proj vs. ctrl.	orientation	—	0.85
7–14	Proj./V+Proj vs. ctrl.	per-relation (4×2)	—	0.25–1.00
15	Two-stage A-lin. vs. ctrl.	agreement (127 boxed)	45/16	0.0003
16	Two-stage A-MLP vs. ctrl.	agreement (127 boxed)	47/14	<0.0001
17	Two-stage B-lin. vs. ctrl.	agreement (127 boxed)	44/20	0.0037
18	Two-stage B-MLP vs. ctrl.	agreement (127 boxed)	47/18	0.0004
19	Zero-shot vs. LM-only	consistency 662	117/43	<0.0001
20	HardNeg vs. LM-only	consistency 662	31/41	0.29
21	Projector vs. LM-only	consistency 662	63/48	0.18
22	V+Proj vs. LM-only	consistency 662	65/56	0.47

Table 4. Complete VSR test battery (22 rows). Per-relation cells (IDs 7–14) are enumerated in the committed vision_side_comparison.json: projector facing $p=0.77$, facing-away $p=1.00$, parallel $p=0.25$, perpendicular $p=1.00$; vision+proj facing $p=1.00$, facing-away $p=0.58$, parallel $p=1.00$, perpendicular $p=1.00$. Rows 15–18 are agreement-based: the committed two-stage pipeline scored agreement with the control’s verdicts (not accuracy against ground truth), so these rows are descriptive and are excluded from the confirmatory set and from the Holm family. SITE transfer tests (3: all/primary/secondary, §E) form a separate small family.

227 adjusted p -values are 0.012 (primary), 1.0, 1.0.

228 Single-checkpoint note. Every trained condition is a
229 single checkpoint. All McNemar tests above quantify
230 paired disagreement on the fixed test examples between
231 two conditions; they do not estimate training-run vari-
232 ability, and the single-seed results reported here do not
233 estimate training-run variability. Adapter-rank claims
234 (e.g., the 1–2pp overall differences among trained 7B
235 conditions) should be read in that light.

236 Consistency verdict patterns. For each strict
237 family and condition, the committed verdict
238 table (results/consistency_verdict_table.json,
239 scripts/consistency_verdict_table.py) reports the
240 observed both-True / both-False / complementary
241 counts, the original and flipped True rates, and the
242 verdict rates expected under verdict independence
243 given those base rates:

The zero-shot pattern is dominated by a False-
defaulting response bias: expected consistency under
verdict independence is 30.2%, below the observed
36.9%, so a simple base-rate null does not explain the
facing inconsistency; conversely, training raises consis-
tency while introducing both-True contradictions (0%
→ 24.3% LM-only / 17.5% hard-neg), a disclosed spe-
cialization trade-off. The frontier-scale zero-shot model
shows the same False-defaulting bias (0 both-True;
56.0% both-False; observed complementary 44.0% vs.
35.6% expected under independence), so scale alone
does not alter the coherence failure mode. Full table
for left/right ($n=245$), front/behind ($n=314$), and par-
allel/perp (soft, $n=34$) in the committed JSON.

Consistency definitions. Strict complements (ex-
actly one truth value): left↔right ($n=245$), in front
of↔behind ($n=314$), facing↔facing away ($n=103$).

Table 5. Facing-pair verdict patterns by condition ($n=103$); “exp.” = expected under verdict independence given the observed base rates.

Condition (facing pairs, $n=103$)	both-T	both-F	comp.	orig-T	flip-T	exp. both-F	exp. comp.
7B zero-shot	0	65 (63.1%)	38 (36.9%)	0.20	0.17	66.5%	30.2%
MiMo-V2.5 zero-shot ($n=100$)	0	56 (56.0%)	44 (44.0%)	0.30	0.14	60.2%	35.6%
LM-only LoRA	25 (24.3%)	10 (9.7%)	68 (66.0%)	0.62	0.52	18.0%	49.4%
HardNeg LoRA	18 (17.5%)	5 (4.9%)	80 (77.7%)	0.60	0.52	19.0%	50.2%
Projector LoRA	18 (17.5%)	15 (14.6%)	70 (68.0%)	0.53	0.50	23.5%	50.3%
Vision+Proj LoRA	14 (13.6%)	23 (22.3%)	66 (64.1%)	0.49	0.43	27.1%	49.4%

Soft complement (parallel $\leftrightarrow$ perpendicular, $n=34$): both-False is legitimate; only both-True is a contradiction; both-True rates are reported separately. Self-consistency on a strict pair means opposite verdicts.

Wilson CIs. All per-relation and SITE subset accuracies are reported with Wilson 95% confidence intervals (main text and committed JSONs).

5. SITE Protocol, Keyword Rule, and Distributions

Pre-specification. Protocol frozen 2026-08-09 before any evaluation; example IDs committed in results/site/site_protocol.json; the protocol configuration is recorded in results/site/run_metadata.json. Evaluation-only: no SITE data used for training. Evaluation: Qwen2-VL-7B-Instruct frozen, greedy, official multiple-choice prompt and letter parser, max_new_tokens=16, images resized to ≤ 392 px long side (constant protocol parameter; the installed transformers build ignores max_pixels), SDPA attention.

Subsets. Primary = official category spatial relationship reasoning (1,721 examples; 993 images). Secondary = heuristic orientation keywords (3,272; 2,024 images). Exploratory = official movement prediction & navigation (1,103; 248 images). Overlaps (image): primary $\cap$ secondary = 488, secondary $\cap$ exploratory = 186, primary $\cap$ exploratory = 0; union = 2,591 images, each evaluated once.

Exact orientation-keyword rule (applied, case-insensitive, word-boundary regex, to question text plus non-image option text; src/datasets/site.py):

orientation, oriented, orient,
facing, face, faces, direction,
directional, view, viewpoint,
view angle, angle, turned, rotate,
rotated, rotation, toward, towards,
away from, left, right, front,
behind, in front of, parallel,
perpendicular, clockwise,
counterclockwise

Secondary subset composition ($n=2,024$ images) by

official category: multi-view & cross-image reasoning 866 (raw 37.8%, CAA 8.8%); spatial relationship reasoning 488 (71.3%, 53.5%); object localization & positioning 346 (46.5%, 28.3%); movement prediction & navigation 186 (49.5%, 24.3%); 3D information understanding 91 (27.5%, 0.4%); counting & existence 47 (53.2%, 37.1%). Top sources: exoego4d 316 (36.4%), egoexo4d 288 (34.7%), MMTBench 237 (48.9%), SAT 183 (46.4%), MMIU 133 (28.6%), CLEVR 106 (92.5%), SPEC 102 (50.0%), SpatialEval 100 (53.0%), ThreeDSRBench 95 (61.1%), MMERealWorld 89 (32.6%), SeedBench 63 (58.7%), GQA 41 (75.6%), VStarBench 37, Blink 35, plus 27 smaller sources.

Decision rule (pre-specified). (1) If the primary subset shows the same broad weakness: run the VSR-trained 7B General LoRA on SITE (pre-specified step 2). (2) Else if only the orientation subset shows the effect: narrow the claim to object/direction orientation. (3) Else: report benchmark-specificity. Outcome: branch 2; the claim was narrowed accordingly. The VSR-trained 7B General LoRA adapter was subsequently evaluated on the identical SITE examples as a post-protocol cross-dataset transfer check (no SITE training; predictions in results/site/vsr_lora_predictions.csv, paired statistics from scripts/compare_site_zeroshot_lora.py, canonical table in results/tables/site_validation_table.md):

Entries are raw accuracy / CAA (%). The adapter is statistically neutral overall, does not improve the orientation subset, and significantly degrades the official spatial-relationship category — VSR gains transfer poorly and can induce negative transfer.

Composition-confound analysis (CPU-only, zero-shot predictions). Because the orientation slice is keyword-derived and heterogeneous, we test whether the heuristic flag predicts lower correctness after controlling for task composition (logistic regression over the 2,591 zero-shot predictions; scripts/site_confound_analysis.py $\rightarrow$ results/site/orientation_confound_analysis.json):

Within the official spatial relationship reasoning category, orientation-flagged questions remain descrip-

Table 6. SITE VSR-LoRA transfer check (post-protocol; predictions in results/site/vsr_lora_predictions.csv, canonical table in results/tables/site_validation_table.md). Subsets are defined by the frozen protocol IDs (see note below); no keyword reconstruction is used in analysis.

Subset	Zero-shot	VSR-LoRA	Δ	McNemar p
All images ($n=2,591$)	54.2 / 31.1	53.8 / 30.5	-0.4	0.66
Primary: spatial-rel. ($n=993$)	75.1 / 59.2	71.6 / 53.4	-3.5	0.004
Secondary: orient. kw. ($n=2,024$)	48.3 / 24.0	48.7 / 24.5	+0.4	0.70
Exploratory: movement ($n=248$)	48.8 / 23.1	50.4 / 25.6	+1.6	0.64

Table 7. Composition-confound logistic regressions (2,591 zero-shot SITE predictions; outcome = correct).

Model: correct $\sim$ orient + controls	OR (orient)	95% CI	p
Unadjusted	0.31	[0.25, 0.38]	2.4×10^{-28}
+ official category	0.74	[0.57, 0.96]	0.023
+ category, source, modality, $n_{options}$	0.84	[0.60, 1.16]	0.29

tively 7.5pp harder than unflagged ones (71.3% vs. 78.8%; $n=488/505$). No source dataset has $n \geq 30$ in both arms (orientation-heavy corpora are almost entirely flagged, others almost never), so source indicators absorb most of the orientation variance in the full model. Reading: the adjusted association is in the expected direction but not significant; the low aggregate heuristic score is largely explained by task composition. We therefore do not treat SITE as independent confirmation of the VSR orientation bottleneck, and the main text reports the slice as exploratory.

High-precision post-hoc subset (exploratory only). Restricting the heuristic to VSR-construct terms only (facing, facing away, parallel, perpendicular; word-boundary match on question or option text) yields $n=177$ examples (77 in the official spatial-relationship category, 99 in multi-view/cross-image reasoning, 1 in object localization): raw accuracy 44.1% (95% CI [37.0, 51.4]), CAA -2.3%. This subset is explicitly post-hoc and does not replace the pre-specified frozen definition; it is reported as exploratory only.

6. Failure Taxonomy

The 48 persistent-failure statements are the union of: wrong in both 7B zero-shot and 7B LoRA (20); wrong in ≥ 3 of 4 base conditions (38); wrong in both 2B and 7B zero-shot (28). The taxonomy below was produced by an automated audit (LLM-generated labels; commit 4b83729); its image grounding is unverified, and its reliability is assessed against an independent human re-audit below. The automated audit suggests many persistent errors remain visually interpretable while also identifying viewpoint and annotation ambiguities; it is qualitative evidence, not a causal localization of the failure. Annotations committed in results/failure_annotations.csv and re-

sults/orientation_persistent_annotations.csv.

Failure mode	n	Share
Clear image, model reasoning failure	18	37.5%
Camera viewpoint ambiguity	8	16.7%
Parallel/perpendicular geometry	6	12.5%
Annotation questionable (label wrong)	5	10.4%
Intrinsic orientation ambiguous	4	8.3%
Front/back object ambiguous	4	8.3%
Small / occluded object	2	4.2%
Subject/reference inversion	1	2.1%

Table 8. Persistent-failure taxonomy, 48 annotated cases.

Transition counts. Scaling (2B→7B zero-shot): 28 both wrong, 23 2B-wrong/7B-right, 22 2B-right/7B-wrong, 64 both right. Adaptation (7B zero→7B LoRA): 20 both wrong, 30 fixed, 27 broken, 60 both right. Of the 48 annotated cases, hard-negative training fixes 5; 15 still fail under all three 7B conditions. Facing/facing-away account for 68.8% of persistent failures.

Human re-audit and inter-source agreement (additive reliability evidence). Because the clean-label audit (§3) and the taxonomy above originate from an automated (LLM-generated) audit, an independent human annotator re-rated the same examples from the images under a blind protocol: scripts/export_iaa_sheets.py exports two rating sheets (the 137-example orientation test set for the binary clean/ambiguous flag, and the 48 persistent-failure cases for the eight-class taxonomy) containing only image, statement, and relation—ground truth, model predictions, and the automated audit’s labels are withheld—and scripts/iaa_tool.py provides a blind annotation interface with auto-save. Agreement is computed by scripts/compute_iaa.py (Cohen’s kappa for the binary flag, Krippendorff’s alpha, nominal, for the taxon-




Figure 1. Representative persistent failures (annotated grid; 15 cases). Images: VSR test set. Per the automated audit (corroborated by the human re-audit below), the most frequent class is clear-image reasoning failure: the pose is visually interpretable to a human and the model still errs.

omy, both with bootstrap 95% confidence intervals; results/iaa/iaa_results.json).

Results (human vs. automated audit). On the 48 cases rated by both sources (the persistent-failure set; the remaining 89 orientation examples were never flagged by the automated audit and are excluded from the agreement statistics), the binary clean/ambiguous flag agrees on 32/48 cases (66.7%; $\kappa=0.36$, bootstrap 95% CI [0.13, 0.59]). On the eight-class taxonomy the exact-class agreement is 15/48 (31.2%; $\alpha=0.14$, CI [-0.03, 0.28]). The disagreement is systematic rather than random: the human annotator never used camera-viewpoint ambiguity or subject/reference inversion, over-used annotation questionable (11 vs. 5) and small/occluded object (11 vs. 2), and flagged 62 of the 137 orientation examples as ambiguous, versus 30 excluded by the automated audit’s strict set. Part of the disagreement on annotation questionable is attributable to the blind protocol itself: judging whether a ground-truth label is wrong requires the label, which was withheld.

Consequences. The frozen automated-audit table (Table 2) is retained unchanged; we do

not redefine the exclusion sets to manufacture agreement. The human re-audit table (Table 3) reproduces the qualitative conclusions despite low per-item agreement, so the sensitivity result is robust at the aggregate level while per-item flags are annotator-specific. An annotator-dependence analysis (results/iaa/consensus_subsets.json, scripts/iaa_consensus_subsets.py) shows the same pattern: under the human’s binary masks the clean-subset gains largely vanish (7B General LoRA 66.7% and 7B zero-shot 54.7% on $n=75$), while under the consensus union mask ($n=62$) the trained conditions again improve (7B General LoRA 77.4%, 7B Hard-Neg 74.2%, 2B zero-shot 72.6%) but 7B zero-shot does not (58.1%). The statement that 2B structured prompting shows no clean-subset gain (53.3% on both full and strict sets) is specific to the taxonomy-derived masks: under the human binary consensus mask it reaches 69.3%. Overall the human re-audit provides partial but aggregate-level corroboration, so the clean-label analysis should be read as annotator-specific exploratory evidence rather than a ruling on label noise. No reported accuracy number is altered by this protocol.

LLM audit reliability note. We also ran the audit prompts through a second LLM (MiMo-V2.5, image-grounded; results/mimo/audit_mimo_*.csv). Its labels collapsed to a single dominant class (43/48 taxonomy cases and 132/137 binary flags), yielding no agreement with either the automated first-pass labels or the human re-audit ($\kappa=0.00$; $\alpha \approx -0.1$; results/iaa/human_mimo_results.json, results/iaa/llm_mimo_results.json). We report this as evidence that naive LLM audit prompting can degenerate to the first option, and we do not substitute it for the human re-audit.

7. Reproducibility and Artifact Index

Code: anonymized archive accompanying this submission (no external links). Every headline number in this paper maps to a committed artifact; the audit script scripts/make_paper_figures.py recomputes the main tables directly from the raw prediction CSVs (output archived at paper/fig/audit_output.txt).

Environments. VSR/7B and SITE runs: Ubuntu 24.04, Python 3.12.13, torch 2.12.1+cu130, transformers 5.14.1, datasets 5.0.1, peft 0.20.0, 1× RTX A6000 (48 GB). 2B runs: earlier python3.10-era environment (SmoIVLM2 wrapper). MiMo-V2.5: public API (OpenCode Go gateway), thinking disabled, 12 parallel workers, per-request token accounting in results/mimo/usage_report.json (total spend reported there). No lock files exist; the adapter configs (com-

Headline number	Artifact
74.0 / 76.6 / 76.5 / 68.3 (2B)	smolvlm2_*_metrics_*.json
80.9 / 84.7 / 84.3 / 82.9 / 83.1 (7B)	*_metrics_*.json
62.8 / 63.5 / 65.7 / 66.4 (orient. family)	same + prediction CSVs
Per-relation cells (Table 1 (main))	orientation_analysis.json
48-failure taxonomy	orientation_persistent_*.json
Probe tables (App. 3)	probe/*.json
Two-stage (App. D, agreement-based rows 15–18)	two_stage_results.json
Consistency (Fig. 4 (main))	consistency_stats_all.json
Vision-side p-values	vision_side_comparison.json
SITE numbers	site/zeroshot_7b_predictions.csv
	site/zeroshot_image_metrics.json
MiMo-V2.5 VSR / consistency / audit (Table 1 row, verdict table, App. 3)	mimo/*.csv, mimo/mimo_vsr_summary.json
Human re-audit + inter-source agreement (App. 6)	iaa/rater2_*.csv, iaa/iaa_results.json, iaa/consensus_subsets.json, iaa/

Table 9. Headline-number → artifact → commit index (all paths relative to results/).

478	mitted) pin the LoRA hyperparameters. Seeds: 42	sults/site/site_protocol.json → frozen_ids.secondary	516
479	defaults in 2B paths; single default seed elsewhere —	(3,272 examples; 2,024 with images), intersected	517
480	per-run seeds are not recorded in the metrics JSONs	with the 2,591 evaluated IDs — is the authoritative	518
481	(a known limitation).	subset, and all numbers in this paper use it ($n=2,024$;	519
482	Canonical freeze. The experimental result set	zero-shot 48.3% / 24.0% CAA; VSR-LoRA 48.7% /	520
483	is a frozen snapshot. Canonical tables are gener-	24.5%, $p=0.70$). The derived question-only figure was	521
484	ated by scripts/build_canonical_tables.py; clean-	corrected before submission; both definitions support	522
485	label results by scripts/clean_label_orientation.py	the same conclusion (orientation $\ll$ primary), but only	523
486	(automated-audit table) and	the frozen-ID set matches the pre-specification. No	524
487	scripts/build_human_clean_label_table.py (ver-	prediction CSV was modified and no model was rerun.	525
488	sioned human re-audit table); publication fig-	Data licenses/access. VSR: cam-	526
489	ures by scripts/make_publication_figures.py;	bridgeltl/vsr_random (HF). SITE: franky-	527
490	SITE paired transfer statistics by	veteran/SITE-Bench, CC-BY-4.0 (media in 130 GB	528
491	scripts/compare_site_zeroshot_lora.py;	zips; requires HF read token). LoRA adapters are	529
492	VSR/SITE headline recomputation by	committed and loadable with peft; base-model weights	530
493	scripts/make_paper_figures.py (output archived	via HF Hub.	531
494	at paper/fig/audit_output.txt); MiMo-V2.5 evalua-	Exact artifacts for the probe tables	532
495	tion and analysis by scripts/eval_mimo_vsr.py and	(App. 3): results/probe/probe_results.json,	533
496	scripts/mimo_analysis.py (predictions and summaries	results/probe/patch_probe_results.json, re-	534
497	under results/mimo/).	sults/probe/grounded_probe_results.json, re-	535
498	Two-stage reproducibility caveat. The two-	sults/probe/clear_subset_results.json, re-	536
499	stage pipeline requires per-patch embeddings	sults/probe/grounded_boxes.json.	537
500	(results/probe/patch_embeddings.pkl, 8.2 GB,		
501	not committed; regenerable on GPU via		
502	scripts/extract_patch_embeddings.py) and hard-		
503	codes a Linux working directory. The committed evi-		
504	dence is the aggregate results/two_stage_results.json		
505	(agreement-based, see App. D); the per-example		
506	decisions are not committed.		
507	SITE secondary-subset audit note (corrected		
508	before submission). An earlier derived analysis		
509	(compare_site_zeroshot_lora.py v1, which fed		
510	the initial zeroshot_image_metrics.json-adjacent		
511	figures and the first canonical site table) re-		
512	constructed the secondary subset by applying		
513	ORIENTATION_KEYWORDS to the question text		
514	only, yielding $n=1,824$ (47.3% raw / 22.6% CAA).		
515	The frozen protocol’s pre-specified definition — re-		