

Beyond Spatial Accuracy: Diagnosing Persistent Orientation Failures in Vision–Language Models

Anonymous authors

Paper under double-blind review

Abstract

Aggregate spatial accuracy can hide sharply different capabilities within vision–language models (VLMs). We study object-relative orientation as a diagnostic case. On Visual Spatial Reasoning (VSR), moving from a compact 2.2B model to a larger 7B VLM substantially improves overall, depth, horizontal, and containment accuracy, but orientation changes only modestly. Except for structured prompting, which degrades performance, trained conditions on both models remain in a narrow 62–66% orientation band. We test structured prompting, general and targeted LoRA, hard negatives, visual/projector adaptation, frozen and object-grounded representation probes, and object-centric decomposition; none of the tested configurations reliably closes the gap, and a single-annotator clean-label sensitivity analysis shows that the pattern persists after excluding examples flagged as ambiguous. Accuracy also fails to capture relational coherence. LM-side adaptation significantly improves consistency over zero-shot, while hard-negative training raises facing/facing-away consistency from 66.0% to 77.7% with facing accuracy unchanged at 68.9%; the additional pooled consistency gain over LM-only is not significant. Cross-dataset evaluation on SITE shows that VSR-trained adaptation does not transfer cleanly: it is neutral overall and significantly harms official spatial-relationship reasoning (75.1%→71.6%, $p = 0.004$). A pre-specified orientation-vocabulary slice, frozen before SITE evaluation, is analyzed as exploratory evidence rather than as direct replication of the VSR construct. The results point to a persistent orientation bottleneck in which orientation is not robustly decodable by the tested frozen-feature readouts and relational decisions are not always propagated coherently, while also showing that task-specific spatial adaptation can transfer poorly. Every headline number was recomputed from raw prediction artifacts by an independent audit script; prompts, adapter configurations, complete probe tables, and the statistical battery are provided in the appendices for reproducibility.

1 Introduction

Spatial reasoning is central to multimodal systems that must interact with physical environments. A useful VLM should not only recognize objects, but determine whether one object is in front of another, whether it is facing another, whether two objects are parallel, and whether equivalent or complementary formulations yield logically compatible answers. These abilities matter for robotics, embodied agents, navigation, and scene understanding.

Existing work already shows that spatial relations are difficult for vision–language systems. VSR was designed as a controlled natural-image probe and reported particularly poor performance on object-orientation relations in earlier VLMs Liu et al. (2023). More recent controlled benchmarks likewise find that strong VLMs can fail on basic left/right and other grounded relations Kamath et al. (2023); Rajabi & Kosecka (2024), while broader evaluations continue to identify spatial orientation as a difficult factor Wang et al. (2025). Our question is therefore not whether spatial reasoning can fail, but whether the orientation weakness persists in modern generative VLMs of different sizes and after targeted intervention, and what the failure signature reveals about representation, reasoning, and transfer.

In our VSR experiments, the distinction between relation families is striking. Moving from SmolVLM2-2.2B-Instruct to the larger Qwen2-VL-7B-Instruct increases zero-shot overall accuracy from 74.0% to 80.9%. Horizontal reasoning rises from 70.2% to 84.8%, depth from 68.9% to 75.2%, and containment from 83.4% to 89.3%. Orientation, by contrast, moves only from 62.8% to 63.5%. We then evaluate interventions spanning prompting, language-side adaptation, targeted data construction, visual-side adaptation, representation probing, and explicit object-centric decomposition. Structured prompting significantly hurts performance; general and targeted LoRA improve aggregate accuracy but do not remove the orientation ceiling; hard negatives improve targeted relational coherence without a significant overall orientation-accuracy gain; and the tested visual-side LoRA configurations do not outperform the language-side control.

As a sensitivity analysis, removing examples flagged by a single-annotator audit raises orientation accuracy (65.7%→78.5% for 7B General LoRA, 66.4%→76.6% for hard-negative LoRA on the strictest 107-example subset) but does not eliminate error (2B structured shows no gain). Because the removed examples were drawn from persistent failures, part of this gain is expected by construction.

A second failure mode emerges from logical consistency. The zero-shot 7B model is consistent on only 36.9% of complementary facing/facing-away pairs. LM-side LoRA substantially improves coherence, and hard-negative training raises facing consistency to 77.7% while facing accuracy remains unchanged at 68.9%. The extra hard-negative gain is not significant in the pooled strict-family comparison, yet accuracy stays unchanged while consistency moves: the two metrics capture different behavior.

Finally, we test cross-dataset transfer. The 7B model is strong on SITE’s official spatial-relationship-reasoning category (75.1% raw; 59.2% chance-adjusted accuracy), yet the VSR-trained adapter does not generalize cleanly: it is statistically neutral overall (54.2%→53.8%, $p = 0.66$) and significantly degrades the official category (75.1%→71.6%, $p = 0.004$). A pre-specified, non-official orientation-vocabulary slice is much weaker in the aggregate (48.3% raw; 24.0% chance-adjusted) and shows no positive transfer, but—as we detail below—its composition limits what it can establish, and we therefore do not treat SITE as independent confirmation of the VSR orientation construct.

Our contributions are fourfold:

- **Fine-grained diagnosis across two differently sized VLMs.** Across ten VSR conditions and two differently sized models, several spatial families improve substantially while object-relative orientation remains unusually resistant.
- **Robustness and mechanism analysis.** A single-annotator clean-label sensitivity analysis, frozen/object-grounded probes, visual-side adaptation, and object-centric decomposition constrain purely annotation-based or single-component explanations.
- **VSR-specific accuracy-coherence separation.** Complementary-relation tests on VSR’s strict relation pairs show that, for the tested conditions, relational consistency can change substantially while task accuracy remains fixed.
- **Cross-dataset transfer analysis.** SITE shows strong official spatial-relationship performance, while VSR-trained LoRA transfers poorly and significantly harms the official SITE category; the orientation-vocabulary slice is analyzed as exploratory evidence.

2 Related Work

Spatial reasoning benchmarks. VSR introduced a controlled natural-image benchmark of more than 10k image-text pairs spanning 66 spatial relations and explicitly highlighted orientation relations such as *facing* and *parallel* as difficult for the VLMs studied at the time Liu et al. (2023). Our work should therefore be read as a modern diagnostic extension rather than the first observation of an orientation weakness. What’sUp isolates basic spatial relations by holding object identities fixed while varying configurations, and found large gaps between humans and tested vision-language models Kamath et al. (2023). GSR-Bench extends controlled grounded spatial evaluation across a larger set of multimodal models and model scales Rajabi & Kosecka (2024). SITE broadens evaluation to single-image, multi-image, and video inputs and organizes tasks

using cognitive-science dimensions including spatial visualization/orientation, intrinsic/extrinsic relations, and static/dynamic settings; its authors report spatial orientation as a persistent challenge Wang et al. (2025). Recent broad evaluations such as SpatialScore similarly argue that spatial intelligence remains fragmented across tasks and modalities Wu et al. (2026), and Mind the Gap isolates spatial reasoning from detection and semantic comprehension in contemporary VLMs Stogiannidis et al. (2025).

2026 landscape and our scope. Since this study’s experiments were frozen, several large-scale 2026 evaluations have reinforced the general picture: SpinBench evaluates 43 models on perspective and rotation, reporting egocentric bias and rephrasing inconsistency Yin et al. (2026); OrientSAM identifies object orientation as a bottleneck for reference-centric reasoning and proposes an orientation-aware training fix OrientSAM authors (2026); and Spatial-DISE and OmniSpatial broaden cognitively organized spatial benchmarking at ICLR 2026 Spatial-DISE authors (2026); Jia et al. (2026). These works corroborate the phenomenon at greater breadth but do not run the controlled intervention ladder we report here (adaptation variants, frozen/object-grounded probes, visual-side adaptation, complementary-relation consistency, clean-label sensitivity, and cross-dataset transfer on the same checkpoints). Our contribution is therefore best read as the mechanistic diagnostic complement to the broad-benchmark line: what changes, and what demonstrably does not change, when standard interventions are applied to the orientation family.

Diagnosing spatial failures. Recent work has begun to analyze spatial failures mechanistically rather than only benchmarking them. On VSR itself, AdaptVis traces attention distributions and finds that successful spatial reasoning correlates with attention aligning to actual object locations Chen et al. (2025). CREG characterizes the directional organization of attribution evidence in Qwen2-VL on spatial-relation benchmarks, showing that standard heatmaps often reflect image layout rather than the queried relation Tan et al. (2026). When More Is Less analyzes when injecting spatial cues, commonsense knowledge, or chain-of-thought instructions helps or hurts VSR, with effects varying by model and setup Akasaka & Han (2026). Probing Visual Concepts measures linear decodability of visual concepts, including orientation-relevant ones, in lightweight driving VLMs and identifies information-flow bottlenecks Theodoridis et al. (2026). VSP decomposes visual spatial planning into perception and reasoning subtasks Wu et al. (2025), and MM-Spatial targets 3D spatial understanding with geometric inputs Daxberger et al. (2025).

Prior work has thus already established spatial and orientation weaknesses in VLMs and has begun to analyze those failures through prompting, attention, representation probing, and consistency-based training. Our contribution is therefore not the first identification of orientation difficulty or spatial inconsistency. Instead, we revisit the known VSR orientation weakness in modern generative VLMs and subject the same relation family to a controlled diagnostic ladder spanning adaptation, frozen and object-grounded readouts, visual-side interventions, complementary-relation consistency, label sensitivity, and cross-dataset transfer.

Adaptation and spatial specialization. We use LoRA, which adapts frozen models through trainable low-rank updates Hu et al. (2021), to compare general, targeted, hard-negative, projector, and vision-plus-projector interventions. This design lets us ask whether benchmark gains reflect general spatial competence or specialization. The external SITE transfer experiment is central here: VSR adaptation is neutral overall on SITE, does not improve the orientation subset, and significantly degrades SITE’s official spatial-relationship category. The result complements benchmark work by explicitly measuring negative as well as positive cross-dataset transfer.

Representation probing and logical consistency. Diagnostic probes can reveal whether information is accessible from frozen representations, but probe performance must be interpreted cautiously because probe capacity can itself learn the target task Hewitt & Liang (2019). We therefore use linear and nonlinear probes, compare against majority baselines, add object-grounded features, and avoid claiming that weak decodability proves information is absent. Separately, consistency across transformed or complementary statements has prior literature, including consistency-oriented training objectives such as SAGE’s geometric and linguistic duality operations Liu et al. (2026). Our consistency analysis is a VSR-specific diagnostic rather than a new consistency framework: it uses strict complementary relation pairs with observed verdict-pattern and base-rate analysis, and it shows that instance-level accuracy and complementary-relation coherence can diverge for the tested conditions.

3 Experimental Setup

3.1 Benchmarks

VSR. We use VSR Liu et al. (2023) as the primary diagnostic benchmark. The random-split test set contains 2,195 examples. The benchmark defines 64 relation labels; 61 appear in the evaluated test split. The orientation family contains *facing*, *facing away from*, *parallel to*, and *perpendicular to*, totaling 137 test examples. All main VSR results use the same held-out test set and deterministic True/False parsing.

SITE. We use SITE Wang et al. (2025) for cross-dataset transfer evaluation only and do not train on SITE. The image evaluation contains 2,591 examples: 1,368 single-image and 1,223 multi-image examples. Our primary SITE subset is the official spatial-relationship-reasoning category ($n = 993$). We additionally report a transparently defined, non-official keyword-derived orientation subset ($n = 2,024$), analyzed as exploratory evidence rather than as direct replication of the VSR construct. SITE results use the benchmark’s native multiple-choice format and chance-adjusted accuracy (CAA) in addition to raw accuracy.

3.2 Models and Conditions

Our compact model is SmolVLM2-2.2B-Instruct, and our larger model is Qwen2-VL-7B-Instruct Wang et al. (2024). We evaluate ten canonical VSR conditions: 2B zero-shot, 2B structured prompting, 2B General LoRA, 2B Targeted LoRA, 7B zero-shot, 7B General LoRA, 7B Targeted LoRA, 7B Hard-Neg LoRA, 7B Projector LoRA, and 7B Vision+Projector LoRA. LoRA Hu et al. (2021) is used for all parameter-efficient adaptation conditions.

General LoRA uses a proportional sample across relation families. Targeted LoRA allocates most training examples to weak orientation/depth/horizontal families with rehearsal from the remaining relations. Hard-negative LoRA introduces audited complementary orientation examples, primarily facing/facing-away inversions. Vision-side conditions adapt either the multimodal projector or upper vision blocks together with the projector. The frozen experimental snapshot records the final result set.

3.3 Metrics and Statistical Tests

For VSR we report accuracy overall and by spatial family. For SITE we report raw accuracy and CAA. Paired condition comparisons use McNemar tests when applicable. Clean-label orientation analysis reports the full orientation set and progressively stricter audited subsets. Logical consistency is evaluated by constructing complementary statements for strict relation pairs (left/right, front/behind, facing/facing-away) and testing whether the model assigns opposite verdicts as required. Parallel/perpendicular is treated separately as a soft complement because both relations can be false for oblique objects.

4 Main Results

4.1 Orientation Remains Weak Across Two Differently Sized VLMs

Table 1 shows the central pattern. Moving from the 2B model to the 7B model lifts zero-shot overall accuracy by 6.9 pp, horizontal by 14.6 pp, depth by 6.3 pp, and containment by 5.9 pp, but orientation by only 0.7 pp. General 7B LoRA further raises overall accuracy to 84.7%, yet orientation reaches only 65.7%. Across trained conditions other than structured prompting, orientation remains in a narrow 62–66% band.

Structured prompting is a negative result: overall accuracy falls from 74.0% to 68.3% at 2B, while orientation falls from 62.8% to 53.3%. A paired exact McNemar test on the full test set confirms that the structured prompt is worse than the minimal baseline ($p \approx 2.2 \times 10^{-7}$). Consistent with recent analyses of information injection on VSR, where additional spatial cues or reasoning instructions can help or hurt depending on the model and setup Akasaka & Han (2026), this result is specific to the tested 2B condition and prompt; it does not establish that prompting generally fails.

Table 1: Canonical VSR test accuracy (%). Orientation improves little with model size or adaptation compared with several other relation families.

Condition	Overall	Orientation	Depth	Horizontal	Containment	Topology/contact
2B zero-shot	74.0	62.8	68.9	70.2	83.4	80.5
2B structured	68.3	53.3	64.9	66.7	78.7	71.7
2B General LoRA	76.6	62.0	71.1	74.3	87.6	81.0
2B Targeted LoRA	76.5	64.2	70.8	75.1	87.6	81.4
7B zero-shot	80.9	63.5	75.2	84.8	89.3	80.3
7B General LoRA	84.7	65.7	82.3	87.4	92.9	84.4
7B Targeted LoRA	83.9	64.2	81.7	87.1	91.7	85.3
7B Hard-Neg LoRA	84.3	66.4	80.7	87.1	89.9	84.6
7B Projector LoRA	82.9	64.2	78.6	87.1	89.3	85.3
7B Vision+Projector LoRA	83.1	64.2	78.0	87.1	89.9	84.4




Figure 1: Overall versus orientation performance across the two VLMs and adaptation conditions.

4.2 Clean-Label Sensitivity Analysis

A single-annotator exploratory audit identified ambiguous or questionable orientation examples, so we repeat evaluation on progressively stricter subsets as a sensitivity analysis (full table in the supplementary). The best full-set orientation result is 66.4%; on the strictest 107-example subset, the best result rises to 78.5% for 7B General LoRA. Because auditing changes the evaluated sample, these values should not be read as a matched comparison against other relation families. The robust conclusion is narrower: removing flagged examples raises orientation scores but does not eliminate error across conditions; the audit is exploratory qualitative evidence, not a ruling on label noise.

5 Why Does Orientation Remain Difficult?

5.1 Frozen Representation Probes

We probe frozen Qwen2-VL-7B-Instruct visual representations without fine-tuning the generative model. Mean-pooled ViT and post-merger features support only weak relation decoding for facing versus facing-away, parallel versus perpendicular, and a four-way orientation task. Linear and MLP probes generally remain near majority baselines; preserving patch-level structure yields a suggestive point estimate for parallel/perpendicular (patch-vote 73.5% vs. 64.7% majority), but that test set is small ($n=34$) and the confidence interval overlaps the majority baseline. In short: orientation is not robustly decodable by the tested frozen-feature readouts.




Figure 2: Orientation performance across 7B interventions, including clean-label analysis and per-relation behavior.

To reduce the confound that global features might infer relations without identifying the relevant objects, we additionally extract subject/reference regions and construct grounded visual features from subject, object, difference, and elementwise-product representations. Object-box conditioning with the model’s own grounding predictions and mean-pooled region features does not recover robust orientation decodability. Following standard cautions about diagnostic classifiers Hewitt & Liang (2019), we phrase these results as decodability under the tested readouts, not as proof that the information is absent from the network. Linear decodability of visual concepts has been probed in driving-oriented VLMs Theodoridis et al. (2026), and directional information in this model family has been inspected via attribution Tan et al. (2026); our probe analysis is not the first of either kind. It differs by targeting VSR’s orientation relation family with pooled, patch-vote, and object-grounded readouts over the exact conditioning features of the generative model, and by pairing decodability with adaptation and consistency interventions rather than treating it as an end in itself. Detailed probe tables and plots are in the supplementary.

5.2 Visual-Side Adaptation and Object-Centric Decomposition

If the frozen representation were the sole limiting factor, directly adapting visual or projector components might recover orientation performance. For the trained checkpoints evaluated here, it did not: projector LoRA and vision+projector LoRA both produced 64.2% orientation accuracy, below the 65.7% LM-only control, while overall accuracy was nominally lower (82.9% and 83.1% vs. 84.7%; unadjusted $p=0.0043/0.0122$; these effects do not survive a blanket multiple-comparison adjustment and are not load-bearing, App. D). This suggests that the tested low-rank adaptation of the visual pathway does not unlock the ceiling.

We also test a box-based two-stage decomposition. Its four-way relation module classifies facing/facing-away/parallel/perpendicular against ground-truth relation labels at 44–46% cross-validation accuracy (49.9% majority baseline; uniform four-class chance 25%), so the tested relation classifier does not outperform the empirical majority baseline. Consistent with that, the pipeline’s True/False verdicts agree with the end-to-end control on only 55.1–58.3% of the 127 boxed orientation statements; in the committed analysis (10 no-box statements scored wrong), verdicts diverge on 61–65 statements, and the control is correct while the pipeline is not in 44–47 of them. The committed pipeline scored agreement with the control rather than ground-truth accuracy, so we report the tested decomposition as failing to reproduce the generative model’s behavior, not as a direct accuracy comparison. Two-dimensional box geometry is structurally insufficient for intrinsic facing direction, and the tested region representations do not supply enough signal to close the gap.

5.3 Logical Consistency Reveals a Distinct Failure Mode

We construct complementary statements for strict relation families. A coherent model should return opposite truth values for left/right, front/behind, and facing/facing-away formulations of the same object pair. Complementary-pair and transformation-based consistency evaluation has prior literature, including




Figure 3: Self-consistency on complementary facing/facing-away relations across 7B conditions.

consistency-oriented training Liu et al. (2026); our contribution is the VSR-specific diagnostic below, built on strict complementary relation pairs with observed verdict-pattern and base-rate analysis. The zero-shot 7B model is consistent on only 58.0% of left/right pairs, 57.0% of front/behind pairs, and 36.9% of facing pairs. The observed verdict patterns show what drives this: the model affirms either member of a facing pair only 17–20% of the time, a False-defaulting bias that yields both-False verdicts on 63.1% of pairs and no both-True verdicts (0/103); expected consistency under verdict independence given these base rates is 30.2%, below the observed 36.9% (verdict table in the supplementary): the dominant phenomenon is a refusal to affirm either proposition, rather than a preference for contradictory affirmations.

LM-only LoRA raises facing consistency to 66.0% and front/behind consistency to 70.7%; pooled strict-family consistency is significantly better than zero-shot ($p < 0.0001$). Hard-negative training further raises facing consistency to 77.7%. Facing accuracy, however, is identical for LM-only and hard-negative LoRA at 68.9%. The additional hard-negative improvement is not significant in the pooled strict-family McNemar comparison ($p = 0.29$). Paired hard-negative supervision therefore provides a targeted intervention on complementary-relation coherence: it descriptively raises facing-pair consistency while leaving original facing accuracy unchanged, although the additional pooled consistency effect over LM-only is not significant. We do not claim a general coherence algorithm; the key result is separability: identical facing accuracy coexists with substantially different facing consistency. Training also changes the failure mode rather than eliminating it: both-False facing verdicts drop from 63.1% to 9.7% (LM-only) while both-True contradictions rise from 0% to 24.3%, a disclosed specialization trade-off.

6 Cross-Dataset Transfer to SITE

We evaluate the frozen 7B base model and the VSR-trained 7B General LoRA adapter on the same 2,591 SITE image examples using SITE’s native multiple-choice protocol Wang et al. (2025), with no SITE training. Because the pre-specified decision rule resolved to branch 2 (the official category was not broadly weak), this adapter run is a post-protocol cross-dataset transfer check rather than a pre-specified step; the protocol and its outcome are documented in the supplementary. The model is strong on the official spatial-relationship-reasoning category: 75.1% raw accuracy and 59.2% CAA. Applying the VSR-trained adapter does not transfer cleanly: it is statistically neutral overall (54.2%→53.8%, $p = 0.66$) and *significantly degrades* the official category (75.1%→71.6%, $p = 0.004$; 52 examples fixed vs. 87 broken). VSR gains therefore transfer poorly and can induce negative transfer on the independent benchmark; we do not equate benchmark-specific fine-tuning gains with general spatial-reasoning improvement.

The pre-specified, non-official orientation-vocabulary slice is substantially weaker in the aggregate (48.3% raw / 24.0% CAA vs. 75.1% / 59.2% for the official category, and below the 54.2% / 31.1% all-images average) and shows no positive transfer under the adapter (48.3%→48.7%, $p = 0.70$). Because the slice is keyword-

Table 2: SITE cross-dataset transfer (images). Entries are raw accuracy / CAA (%).

Subset	Zero-shot	VSR-LoRA	Δ
All images	54.2 / 31.1	53.8 / 30.5	-0.4
Official spatial-rel.	75.1 / 59.2	71.6 / 53.4	-3.5
Orient. heuristic	48.3 / 24.0	48.7 / 24.5	+0.4



Figure 4: SITE CAA for zero-shot and VSR-trained LoRA. The official spatial-relationship category degrades significantly under transfer.

derived and heterogeneous (43% of its examples come from multi-view/cross-image reasoning, only 24% from the spatial-relationship category), we treat it as exploratory evidence rather than as direct replication of the VSR construct. A composition-confound analysis of the 2,591 zero-shot predictions (supplementary §E) supports this caution: after controlling for official category, source dataset, modality, and number of options, the orientation flag is not significantly associated with correctness (odds ratio 0.84, 95% CI [0.60, 1.16], $p = 0.29$), although it remains descriptively 7.5pp lower within the official spatial-relationship category itself (71.3% vs. 78.8%, $n=488/505$). A post-hoc high-precision subset restricted to VSR-construct terms (facing, facing away, parallel, perpendicular; $n=177$) yields 44.1% raw (95% CI [37.0, 51.4]; CAA -2.3%) and is reported as exploratory only. We therefore do not treat SITE as independent confirmation of the VSR orientation bottleneck; the transfer results stand on the official category, where the model is strong and the adapter demonstrably does harm.

7 Discussion

Orientation is not spatial reasoning broadly. The strongest conclusion is deliberately narrow. The 7B model is strong on several VSR families and SITE’s official spatial-relationship category. The persistent weakness concerns the object-relative orientation family in VSR; on SITE, the orientation-vocabulary slice is analyzed as exploratory evidence rather than as independent confirmation. This framing is consistent with the original VSR observation that orientation relations are difficult Liu et al. (2023), while extending it to modern generative VLMs and a much broader intervention analysis.

Decodability and coherence both matter, as diagnostics. The probes show that orientation is not robustly decodable under the tested frozen-feature readouts, while consistency experiments expose large logical failures even for relations answered above chance. Neither result alone identifies a complete causal mechanism, and both are diagnostic rather than causal-localization claims. Together they support a mixed account in which orientation information is difficult to access reliably under the tested readouts and relational decisions are not always propagated coherently across complementary formulations.

Task adaptation can specialize rather than generalize. General LoRA improves VSR overall accuracy substantially, but its SITE transfer is neutral overall and significantly negative on the official spatial-reasoning subset. Hard negatives improve a targeted coherence behavior while producing relation-specific tradeoffs. Benchmark-specific fine-tuning gains should therefore not be equated automatically with general spatial-reasoning improvement.

8 Limitations

Our main diagnostic test contains only 137 orientation examples, and individual relation subsets such as perpendicular are smaller still ($n=12$); perpendicular results are therefore reported descriptively throughout. The single-annotator clean-label audit reduces ambiguity but also decreases sample size, and it is exploratory qualitative evidence rather than a ruling on label noise. The SITE orientation slice is keyword-derived rather than an official benchmark category; a composition-confound analysis shows its aggregate score is largely explained by task composition, so we do not treat SITE as independent confirmation of the VSR orientation bottleneck. Our probing methods test accessible decodability under specific linear/MLP and pooling choices; failure to decode does not prove information is absent. Visual-side experiments use LoRA rather than full end-to-end retraining. We study two differently sized VLMs from different families; this is a cross-model comparison, not a controlled scale experiment. All trained conditions are single-checkpoint runs: the paired McNemar tests quantify disagreement on the fixed test examples, not training-run variability, and these single-seed results do not estimate it. The SITE evaluation is image-only; SITE video reasoning remains future work. Finally, several interventions are evaluated as negative results, so lack of improvement should be interpreted for the tested configurations rather than as impossibility claims about all prompting, adaptation, or architectural methods.

9 Conclusion

Aggregate spatial accuracy hides a persistent orientation bottleneck in the VLMs studied here. A larger VLM and standard adaptation substantially improve several spatial families while leaving object-relative orientation comparatively resistant. Auditing, probing, visual-side adaptation, and decomposition constrain simple explanations, while complementary-relation tests show that accuracy and coherence are distinct axes of spatial competence. Cross-dataset evaluation on SITE further shows that VSR-specific adaptation transfers poorly and can negatively transfer to official spatial-relationship reasoning. The results motivate spatial evaluation that reports relation-specific performance, transfer, and logical coherence rather than relying on aggregate accuracy alone.

10 Broader Impact

This work is diagnostic evaluation of published models on public benchmarks; it introduces no new data collection involving human subjects and no deployment system. Potential positive impact: finer-grained spatial evaluation can reduce overclaiming about multimodal model capabilities relevant to robotics and navigation. Potential negative impact is limited to the usual risks of benchmark-driven research (over-fitting evaluation practices); we mitigate by reporting per-family results, exploratory analyses as exploratory, and negative results for our own interventions.

References

- Muku Akasaka and Soyeon Caren Han. When more is less: A systematic analysis of spatial and commonsense information for visual spatial reasoning. *arXiv preprint arXiv:2602.21619*, 2026.
- Shiqi Chen, Tongyao Zhu, Ruochen Zhou, Jinghan Zhang, Siyang Gao, Juan Carlos Niebles, Mor Geva, Junxian He, Jiajun Wu, and Manling Li. Why is spatial reasoning hard for vlms? an attention mechanism perspective on focus areas. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. URL <https://arxiv.org/abs/2503.01773>.

- Erik Daxberger, Nina Wenzel, David Griffiths, Haiming Gang, Justin Lazarow, Gefen Kohavi, Kai Kang, Marcin Eichner, Yinfei Yang, Afshin Dehghan, and Peter Gräsch. MM-Spatial: Exploring 3d spatial understanding in multimodal LLMs. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 7395–7408, 2025.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pp. 2733–2743. Association for Computational Linguistics, 2019. doi: 10.18653/v1/D19-1275.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- Mengdi Jia, Zekun Qi, Shaochen Zhang, Wenyao Zhang, XinQiang Yu, Jiawei He, He Wang, and Li Yi. Omnispatial: Towards comprehensive spatial reasoning benchmark for vision language models. In *International Conference on Learning Representations (ICLR)*, 2026.
- Amita Kamath, Jack Hessel, and Kai-Wei Chang. What’s “up” with vision-language models? investigating their struggle with spatial reasoning. *arXiv preprint arXiv:2310.19785*, 2023.
- Fangyu Liu, Guy Emerson, and Nigel Collier. Visual spatial reasoning. *Transactions of the Association for Computational Linguistics*, 11:635–651, 2023.
- Junming Liu, Yuqi Li, Yifei Sun, Maonan Wang, Piotr Koniusz, Yirong Chen, and Ding Wang. Self-evolving spatial reasoning in vision language models via geometric logic consistency. *arXiv preprint arXiv:2605.18162*, 2026.
- OrientSAM authors. Orientsam: Mitigating camera-centric shortcut in multimodal spatial reasoning via orientation-aware spatial alignment. *arXiv preprint arXiv:2607.17657*, 2026.
- Navid Rajabi and Jana Kosecka. GSR-BENCH: A benchmark for grounded spatial reasoning evaluation via multimodal LLMs. *arXiv preprint arXiv:2406.13246*, 2024.
- Spatial-DISE authors. Spatial-dise: A unified benchmark for evaluating spatial reasoning in vision-language models. In *International Conference on Learning Representations (ICLR)*, 2026.
- Ilias Stogiannidis, Steven McDonagh, and Sotirios A. Tsaftaris. Mind the gap: Benchmarking spatial reasoning in vision-language models. *arXiv preprint arXiv:2503.19707*, 2025.
- Kaizhen Tan, Yang Feng, and Heqing Du. CREG: Compass relational evidence graph for characterizing directional structure in VLM spatial-reasoning attribution. *arXiv preprint arXiv:2603.20475*, 2026.
- Nikos Theodoridis, Reenu Mohandas, Ganesh Sistu, Anthony Scanlan, Ciarán Eising, and Tim Brophy. Probing visual concepts in lightweight vision-language models for automated driving. *Transactions on Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2603.06054>.
- Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayiheng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-VL: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024.
- Wenqi Wang, Reuben Tan, Pengyue Zhu, Jianwei Yang, Zhengyuan Yang, Lijuan Wang, Andrey Kolobov, Jianfeng Gao, and Boqing Gong. SITE: towards spatial intelligence thorough evaluation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 9058–9069, 2025.
- Haoning Wu, Xiao Huang, Yaohui Chen, Ya Zhang, Yanfeng Wang, and Weidi Xie. Spatialscore: Towards comprehensive evaluation for spatial intelligence. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 31029–31041, 2026.

Qiucheng Wu, Handong Zhao, Michael Saxon, Trung Bui, William Yang Wang, Yang Zhang, and Shiyu Chang. VSP: Diagnosing the dual challenges of perception and reasoning in spatial planning tasks for MLLMs. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 2270–2280, 2025.

Zixuan Yin et al. Spinbench: Perspective and rotation as a lens on spatial reasoning in vlms. *arXiv preprint arXiv:2509.25390*, 2026.

A Exact Prompts, Decoding, and Parsing

VSR plain prompt (identical for 2B and 7B, all conditions; ‘{statement}’ is the VSR caption):

Look at the image carefully.

Statement: "{statement}"

Is this statement true or false?

Answer with exactly one word:

True or False.

2B structured-decomposition prompt (2B only; single run):

Look carefully at the image.

Statement:

"{statement}"

Determine whether the statement is true
by doing the following:

1. Identify the subject object.
2. Identify the reference object.
3. Determine the position and orientation of the subject relative to the reference object.
4. Pay careful attention to left/right, front/back, viewpoint, and orientation.
5. Decide whether the stated spatial relationship matches the image.

Answer with exactly one word:

True or False.

SITE multiple-choice prompt (replicated from the official lmms-eval task sitebench/utils.py):

image examples (one <image> token per visual):

<image>...

Question: {question}

Options:

A: {opt0}

B: {opt1}

...

Give me the answer letter directly.

The best answer is:

```
# video examples (pre-prompt):
Select the best answer to the following
multiple-choice question based on the
video. Respond with only the letter
of the correct option.
Question: {question}
Options:
A: ...
B: ...
...
Give me the answer letter directly.
The best answer is:
```

Decoding. All generations are greedy: `do_sample=False`, temperature 0, top- p 1.0. `max_new_tokens`: 5 for the 2B wrapper, 128 for 7B VSR runs, 16 for SITE (change validated parse-identical on 125/125 held-out outputs before adoption; recorded in the SITE run metadata).

Parsing. VSR outputs are parsed with an anchored case-insensitive regex (`src/evaluation/parser.py`): strips prefixes ('the answer is', 'answer:', 'response:', 'output:') and trailing punctuation; accepts exact 'true/yes/correct' or 'false/no/incorrect' forms; substring fallback only when exactly one of 'true'/'false' appears. Invalid outputs are counted as wrong, never guessed; invalid rate is 0 on all reported runs. SITE outputs use the official `parse_multi_choice_response` (letter); unparseable responses are counted incorrect and reported.

B Training and Adapter Configurations

All LoRA conditions share: $r=8$, $\alpha=16$, dropout 0.05, AdamW lr 10^{-4} , weight decay 0.01, linear warmup 10% of steps, 2 epochs, bf16, gradient checkpointing, gradient clipping 1.0, vision tower frozen. Adapter configs are committed at `checkpoints/*/final/adapter_config.json`.

Condition	LoRA targets	Trainable	Batch	Manifest
2B General	LM q/k/v/o	~2.3M	microbatch 4×4	general (2,000)
2B Targeted	LM q/k/v/o	~2.3M	same	targeted (2,000)
7B LM-only	LM q/k/v/o	~0.3M	1	general (2,000)
7B HardNeg	LM q/k/v/o	~0.3M	1	hard-negative (2,000)
7B Projector	<code>visual_merger.mlp.0/2</code>	152K	1	general (2,000)
7B Vision+Proj	<code>merger + blocks 24-31</code>	1.46M	1	general (2,000)

Table 3: LoRA conditions and their target modules.

Manifest construction. *General*: 2,000 statements sampled from the VSR train split. *Targeted*: orientation-over-sampled subset of the same split. *Hard-negative*: the orientation block of the general manifest is replaced by audited-clean originals plus paired hard negatives (facing↔facing away from; parallel↔perpendicular) with flipped labels; every included example was audited for label correctness (451 train orientation examples audited: 402 clean, 28 excluded, 21 original-only). 2B runs split 5% of the manifest for validation (seed 42). All manifests derive exclusively from the VSR *train* split; all evaluations use *test* only.

C Probe Methodology and Full Tables

Features. Qwen2-VL-7B-Instruct, vision tower frozen. ViT patch embeddings (1280-d) and post-merger features (3584-d), mean-pooled per image (ungrounded) or per grounding box (grounded). Grounding boxes

from frozen Qwen2-VL grounding (94.6% of 647 orientation images yield both boxes; failures are mostly genuinely absent objects), rescaled to pixels; region features mean-pooled over patch centers inside each box.

Tasks. T1 facing vs. facing away (test $n=103$, majority 64.2%); T2 parallel vs. perpendicular ($n=34$, majority 64.7%); T3 4-way ($n=137$, majority 49.6%). Training on audited-clean VSR train; 5-fold CV; no test examples in training.

Readouts. Linear logistic regression; MLP (256 hidden, ReLU, Adam, early stopping on CV); patch-vote (per-patch linear probe, majority vote over patches); grounded variants with feature sets **visual** ([subj, ref, subj − ref, subj · ref]), **geometry** (12-d normalized box features), and **visual+geometry**.

Ungrounded results (test acc %; balanced acc in parentheses; Wilson 95% CI; majority baselines: T1 64.2%, T2 64.7%, T3 49.6%):

T1 facing vs. away ($n=103$)	Test	Bal. acc	95% CI
ViT linear	65.0	60.4	[55, 74]
ViT MLP	55.3	52.0	[45, 65]
Merger linear	68.9	64.5	[59, 77]
Merger MLP	52.4	49.2	[42, 62]
ViT patch-vote	62.1	—	[53, 71]
Merger patch-vote	60.2	—	[51, 69]

T2 parallel vs. perpendicular ($n=34$)	Test	Bal. acc	95% CI
ViT linear	61.8	61.0	[44, 77]
ViT MLP	67.6	69.3	[50, 82]
Merger linear	64.7	61.4	[47, 79]
Merger MLP	64.7	63.3	[47, 79]
ViT patch-vote	73.5	—	[57, 85]
Merger patch-vote	67.6	—	[51, 81]

T3 4-way ($n=137$)	Test	Bal. acc	95% CI
ViT linear	52.6	40.5	[44, 61]
ViT MLP	52.6	42.9	[44, 61]
Merger linear	53.3	44.0	[45, 61]
Merger MLP	47.4	43.2	[39, 56]
ViT patch-vote	51.8	—	[44, 60]
Merger patch-vote	52.6	—	[44, 61]

Grounded results (CV acc / test acc; majority T1 63.7%, T2 57.0%; visual = [subj, ref, subj − ref, subj · ref]; geometry = 12-d box features; T3 4-way is at chance for every readout, test acc 41.7–52.8% across all 12 cells vs. 49.9% majority, committed JSON):

Readout	T1 visual	T1 geom.	T2 visual	T2 geom.
ViT linear	59.8/61.6	61.1/60.6	51.2/53.6	44.2/50.0
ViT MLP	59.5/66.7	61.8/62.6	45.4/46.4	45.4/46.4
Merger linear	57.2/65.7	61.1/60.6	51.2/53.6	44.2/50.0
Merger MLP	56.3/52.5	61.1/71.7	49.0/53.6	50.0/42.9

The single above-majority cell (geometry-only merger MLP on T1, 71.7% test) has CV 61.1% below the 63.7% majority — not robust — and derives from the grounding model’s own box statistics, not vision-feature content. In short, object-box conditioning with the model’s own grounding predictions and mean-pooled region features does not recover robust orientation decodability under the tested readouts. The T2 patch-vote result (73.5%) is a suggestive point estimate with a wide CI overlapping the majority baseline ($n=34$).

Clear-subset control (canonical frozen snapshot). The 48-example manual audit identified ambiguous or questionable orientation labels. Re-evaluation on progressively stricter subsets (values % from results/tables/orientation_clean_label_table.md, generated by scripts/clean_label_orientation.py): Because auditing changes the evaluated sample, strict-subset

Table 4: Clean-label orientation robustness (canonical frozen snapshot; values % from `results/tables/orientation_clean_label_table.md`).

Condition	Full (137)	−Questionable (132)	Clear (124)	Strict (107)
2B zero-shot	62.8	65.2	66.9	75.7
2B structured	53.3	52.3	51.6	53.3
2B General LoRA	62.0	63.6	64.5	72.9
2B Targeted LoRA	64.2	65.9	68.5	76.6
7B zero-shot	63.5	65.2	68.5	72.0
7B General LoRA	65.7	67.4	70.2	78.5
7B Targeted LoRA	64.2	65.9	69.4	72.0
7B Hard-Neg LoRA	66.4	68.2	70.2	76.6
7B Projector LoRA	64.2	65.9	69.4	73.8
7B Vision+Projector LoRA	64.2	65.2	68.5	73.8

values are not a matched comparison against other relation families. The earlier probe-side clear/strict check (pre-freeze) additionally suggested that removing audit-flagged examples does not change the probe conclusion: the best ungrounded probe moved only 68.9% $\rightarrow$ 71.1% $\rightarrow$ 71.6% (full/clear/strict, $n = 103/97/88$) and the generative T1 readouts 64.1/66.0/68.2 (zero-shot) and 68.9/72.2/78.4 (General LoRA); committed in `results/probe/clear_subset_results.json`.

D Complete Statistical Test Battery

All pairwise tests are exact McNemar tests (binomial test on discordant pairs, two-sided). Discordant counts are reported as (condition-losses / condition-gains) relative to the control named in each row. No test was run and withheld; the battery below is complete for this paper.

ID	Comparison	Subset	Discord.	p
1	HardNeg vs. Gen.	overall (2195)	52/60	0.508
2	HardNeg vs. Gen.	weak families	20/26	0.461
3	Projector vs. ctrl.	overall	—	0.0043
4	V+Proj vs. ctrl.	overall	—	0.0122
5	Projector vs. ctrl.	orientation	—	0.86
6	V+Proj vs. ctrl.	orientation	—	0.85
7–14	Proj./V+Proj vs. ctrl.	per-relation (4×2)	—	0.25–1.00
15	Two-stage A-lin. vs. ctrl.	agreement (127 boxed)	45/16	0.0003
16	Two-stage A-MLP vs. ctrl.	agreement (127 boxed)	47/14	<0.0001
17	Two-stage B-lin. vs. ctrl.	agreement (127 boxed)	44/20	0.0037
18	Two-stage B-MLP vs. ctrl.	agreement (127 boxed)	47/18	0.0004
19	Zero-shot vs. LM-only	consistency 662	117/43	<0.0001
20	HardNeg vs. LM-only	consistency 662	31/41	0.29
21	Projector vs. LM-only	consistency 662	63/48	0.18
22	V+Proj vs. LM-only	consistency 662	65/56	0.47

Table 5: Complete VSR test battery (22 rows). Per-relation cells (IDs 7–14) are enumerated in the committed `vision_side_comparison.json`: projector facing $p=0.77$, facing-away $p=1.00$, parallel $p=0.25$, perpendicular $p=1.00$; vision+proj facing $p=1.00$, facing-away $p=0.58$, parallel $p=1.00$, perpendicular $p=1.00$. Rows 15–18 are *agreement-based*: the committed two-stage pipeline scored agreement with the control’s verdicts (not accuracy against ground truth), so these rows are descriptive and are excluded from the confirmatory set and from the Holm family. SITE transfer tests (3: all/primary/secondary, §E) form a separate small family.

Principal vs. exploratory comparisons. The battery above is fully enumerated; no test was run and withheld. The *principal* (confirmatory) comparisons are: zero-shot vs. LM-only consistency (ID 19) and the SITE transfer tests (§E). All other comparisons — vision-side conditions and their per-relation cells, hard-negative training vs. control, the remaining consistency comparisons, and the two-stage agreement rows (IDs 15–18) — are *exploratory diagnostics* reported with exact unadjusted p -values. Under a Holm adjustment over the 18 accuracy-valid VSR tests ($\alpha=0.05$), only the zero-shot-vs-LM-only consistency test survives (Holm

$p=0.0018$, using the conservative $p<0.0001$ bound); the vision-side overall effects (raw $p=0.0043$, $p=0.0122$; Holm 0.073, 0.195) and all other comparisons do not survive a blanket adjustment and are not load-bearing. For the 3 SITE transfer tests, Holm-adjusted p -values are 0.012 (primary), 1.0, 1.0.

Single-checkpoint note. Every trained condition is a single checkpoint. All McNemar tests above quantify paired disagreement on the fixed test examples between two conditions; they do not estimate training-run variability, and the single-seed results reported here do not estimate training-run variability. Adapter-rank claims (e.g., the 1–2pp overall differences among trained 7B conditions) should be read in that light.

Consistency verdict patterns. For each strict family and condition, the committed verdict table (`results/consistency_verdict_table.json`, `scripts/consistency_verdict_table.py`) reports the observed both-True / both-False / complementary counts, the original and flipped True rates, and the verdict rates expected under verdict independence given those base rates:

Table 6: Facing-pair verdict patterns by condition ($n=103$); “exp.” = expected under verdict independence given the observed base rates.

Condition (facing pairs, $n=103$)	both-T	both-F	comp.	orig-T	flip-T	exp. both-F	exp. comp.
7B zero-shot	0	65 (63.1%)	38 (36.9%)	0.20	0.17	66.5%	30.2%
LM-only LoRA	25 (24.3%)	10 (9.7%)	68 (66.0%)	0.62	0.52	18.0%	49.4%
HardNeg LoRA	18 (17.5%)	5 (4.9%)	80 (77.7%)	0.60	0.52	19.0%	50.2%
Projector LoRA	18 (17.5%)	15 (14.6%)	70 (68.0%)	0.53	0.50	23.5%	50.3%
Vision+Proj LoRA	14 (13.6%)	23 (22.3%)	66 (64.1%)	0.49	0.43	27.1%	49.4%

The zero-shot pattern is dominated by a False-defaulting response bias: expected consistency under verdict independence is 30.2%, below the observed 36.9%, so a simple base-rate null does not explain the facing inconsistency; conversely, training raises consistency while introducing both-True contradictions (0% $\rightarrow$ 24.3% LM-only / 17.5% hard-neg), a disclosed specialization trade-off. Full table for left/right ($n=245$), front/behind ($n=314$), and parallel/perp (soft, $n=34$) in the committed JSON.

Consistency definitions. Strict complements (exactly one truth value): left $\leftrightarrow$ right ($n=245$), in front of $\leftrightarrow$ behind ($n=314$), facing $\leftrightarrow$ facing away ($n=103$). Soft complement (parallel $\leftrightarrow$ perpendicular, $n=34$): both-False is legitimate; only both-True is a contradiction; both-True rates are reported separately. Self-consistency on a strict pair means opposite verdicts.

Wilson CIs. All per-relation and SITE subset accuracies are reported with Wilson 95% confidence intervals (main text and committed JSONs).

E SITE Protocol, Keyword Rule, and Distributions

Pre-specification. Protocol frozen 2026-08-09 before any evaluation; example IDs committed in `results/site/site_protocol.json`; the protocol configuration is recorded in `results/site/run_metadata.json`. Evaluation-only: no SITE data used for training. Evaluation: Qwen2-VL-7B- Instruct frozen, greedy, official multiple-choice prompt and letter parser, `max_new_tokens=16`, images resized to ≤ 392 px long side (constant protocol parameter; the installed transformers build ignores `max_pixels`), SDPA attention.

Subsets. Primary = official category *spatial relationship reasoning* (1,721 examples; 993 images). Secondary = heuristic orientation keywords (3,272; 2,024 images). Exploratory = official *movement prediction & navigation* (1,103; 248 images). Overlaps (image): primary $\cap$ secondary = 488, secondary $\cap$ exploratory = 186, primary $\cap$ exploratory = 0; union = 2,591 images, each evaluated once.

Exact orientation-keyword rule (applied, case-insensitive, word- boundary regex, to question text plus non-image option text; `src/datasets/site.py`):

```
orientation, oriented, orient,
facing, face, faces, direction,
directional, view, viewpoint,
```

view angle, angle, turned, rotate,
rotated, rotation, toward, towards,
away from, left, right, front,
behind, in front of, parallel,
perpendicular, clockwise,
counterclockwise

Secondary subset composition ($n=2,024$ images) by official category: multi-view & cross-image reasoning 866 (raw 37.8%, CAA 8.8%); spatial relationship reasoning 488 (71.3%, 53.5%); object localization & positioning 346 (46.5%, 28.3%); movement prediction & navigation 186 (49.5%, 24.3%); 3D information understanding 91 (27.5%, 0.4%); counting & existence 47 (53.2%, 37.1%). Top sources: exoego4d 316 (36.4%), egoexo4d 288 (34.7%), MMTBench 237 (48.9%), SAT 183 (46.4%), MMIU 133 (28.6%), CLEVR 106 (92.5%), SPEC 102 (50.0%), SpatialEval 100 (53.0%), ThreeDSRBench 95 (61.1%), MMERealWorld 89 (32.6%), SeedBench 63 (58.7%), GQA 41 (75.6%), VStarBench 37, Blink 35, plus 27 smaller sources.

Decision rule (pre-specified). (1) If the primary subset shows the same broad weakness: run the VSR-trained 7B General LoRA on SITE (pre-specified step 2). (2) Else if only the orientation subset shows the effect: narrow the claim to object/direction orientation. (3) Else: report benchmark-specificity. Outcome: branch 2; the claim was narrowed accordingly. The VSR-trained 7B General LoRA adapter was subsequently evaluated on the identical SITE examples as a *post-protocol cross-dataset transfer check* (no SITE training; predictions in `results/site/vsr_lora_predictions.csv`, paired statistics from `scripts/compare_site_zeroshot_lora.py`, canonical table in `results/tables/site_validation_table.md`):

Table 7: SITE VSR-LoRA transfer check (post-protocol; predictions in `results/site/vsr_lora_predictions.csv`, canonical table in `results/tables/site_validation_table.md`). Subsets are defined by the frozen protocol IDs (see note below); no keyword reconstruction is used in analysis.

Subset	Zero-shot	VSR-LoRA	Δ	McNemar p
All images ($n=2,591$)	54.2 / 31.1	53.8 / 30.5	-0.4	0.66
Primary: spatial-rel. ($n=993$)	75.1 / 59.2	71.6 / 53.4	-3.5	0.004
Secondary: orient. kw. ($n=2,024$)	48.3 / 24.0	48.7 / 24.5	+0.4	0.70
Exploratory: movement ($n=248$)	48.8 / 23.1	50.4 / 25.6	+1.6	0.64

Entries are raw accuracy / CAA (%). The adapter is statistically neutral overall, does not improve the orientation subset, and *significantly degrades* the official spatial-relationship category — VSR gains transfer poorly and can induce negative transfer.

Composition-confound analysis (CPU-only, zero-shot predictions). Because the orientation slice is keyword-derived and heterogeneous, we test whether the heuristic flag predicts lower correctness after controlling for task composition (logistic regression over the 2,591 zero-shot predictions; `scripts/site_confound_analysis.py` → `results/site/orientation_confound_analysis.json`):

Table 8: Composition-confound logistic regressions (2,591 zero-shot SITE predictions; outcome = correct).

Model: correct $\sim$ orient + controls	OR (orient)	95% CI	p
Unadjusted	0.31	[0.25, 0.38]	2.4×10^{-28}
+ official category	0.74	[0.57, 0.96]	0.023
+ category, source, modality, $n_{options}$	0.84	[0.60, 1.16]	0.29

Within the official *spatial relationship reasoning* category, orientation-flagged questions remain descriptively 7.5pp harder than unflagged ones (71.3% vs. 78.8%; $n=488/505$). No source dataset has $n \geq 30$ in both arms (orientation-heavy corpora are almost entirely flagged, others almost never), so source indicators absorb most of the orientation variance in the full model. **Reading:** the adjusted association is in the expected direction but not significant; the low aggregate heuristic score is largely explained by task composition. We therefore

do not treat SITE as independent confirmation of the VSR orientation bottleneck, and the main text reports the slice as exploratory.

High-precision post-hoc subset (exploratory only). Restricting the heuristic to VSR-construct terms only (facing, facing away, parallel, perpendicular; word-boundary match on question or option text) yields $n=177$ examples (77 in the official spatial-relationship category, 99 in multi-view/cross-image reasoning, 1 in object localization): raw accuracy 44.1% (95% CI [37.0, 51.4]), CAA -2.3% . This subset is explicitly post-hoc and does not replace the pre-specified frozen definition; it is reported as exploratory only.

F Failure Taxonomy

The 48 persistent-failure statements are the union of: wrong in both 7B zero-shot and 7B LoRA (20); wrong in ≥ 3 of 4 base conditions (38); wrong in both 2B and 7B zero-shot (28). A single-annotator exploratory audit (eight-class taxonomy) suggests many persistent errors remain visually interpretable while also identifying viewpoint and annotation ambiguities; it is qualitative evidence, not a causal localization of the failure. Annotations committed in `results/failure_annotations.csv` and `results/orientation_persistent_annotations.csv`.

Failure mode	n	Share
Clear image, model reasoning failure	18	37.5%
Camera viewpoint ambiguity	8	16.7%
Parallel/perpendicular geometry	6	12.5%
Annotation questionable (label wrong)	5	10.4%
Intrinsic orientation ambiguous	4	8.3%
Front/back object ambiguous	4	8.3%
Small / occluded object	2	4.2%
Subject/reference inversion	1	2.1%

Table 9: Persistent-failure taxonomy, 48 annotated cases.

Transition counts. Scaling (2B $\rightarrow$ 7B zero-shot): 28 both wrong, 23 2B-wrong/7B-right, 22 2B-right/7B-wrong, 64 both right. Adaptation (7B zero $\rightarrow$ 7B LoRA): 20 both wrong, 30 fixed, 27 broken, 60 both right. Of the 48 annotated cases, hard-negative training fixes 5; 15 still fail under all three 7B conditions. Facing/facing-away account for 68.8% of persistent failures.

G Reproducibility and Artifact Index

Code: anonymized archive accompanying this submission (no external links). Every headline number in this paper maps to a committed artifact; the audit script `scripts/make_paper_figures.py` recomputes the main tables directly from the raw prediction CSVs (output archived at `paper/fig/audit_output.txt`).

Headline number	Artifact	Commit
74.0 / 76.6 / 76.5 / 68.3 (2B)	<code>smolvlm2*_metrics*.json</code>	—
80.9 / 84.7 / 84.3 / 82.9 / 83.1 (7B)	<code>*_metrics*.json</code>	—
62.8 / 63.5 / 65.7 / 66.4 (orient. family)	same + prediction CSVs	same
Per-relation cells (Table 1 (main))	<code>orientation_analysis.json</code>	—
48-failure taxonomy	<code>orientation_persistent*.json</code>	—
Probe tables (App. C)	<code>probe/*.json</code>	—
Two-stage (App. D, agreement-based rows 15–18)	<code>two_stage_results.json</code>	—
Consistency (Fig. 4 (main))	<code>consistency_stats_all.json</code>	—
Vision-side p-values	<code>vision_side_comparison.json</code>	—
SITE numbers	<code>site/zeroshot_7b_predictions.csv</code>	—
	<code>site/zeroshot_image_metrics.json</code>	—

Table 10: Headline-number $\rightarrow$ artifact $\rightarrow$ commit index (all paths relative to `results/`).

Environments. VSR/7B and SITE runs: Ubuntu 24.04, Python 3.12.13, torch 2.12.1+cu130, transformers 5.14.1, datasets 5.0.1, peft 0.20.0, $1\times$ RTX A6000 (48 GB). 2B runs: earlier python3.10-era environment

(SmolVLM2 wrapper). No lock files exist; the adapter configs (committed) pin the LoRA hyperparameters. Seeds: 42 defaults in 2B paths; single default seed elsewhere — per-run seeds are not recorded in the metrics JSONs (a known limitation).

Canonical freeze. The experimental result set is a frozen snapshot. Canonical tables are generated by `scripts/build_canonical_tables.py`; clean-label results by `scripts/clean_label_orientation.py`; publication figures by `scripts/make_publication_figures.py`; SITE paired transfer statistics by `scripts/compare_site_zeroshot_lora.py`; VSR/SITE headline recomputation by `scripts/make_paper_figures.py` (output archived at `paper/fig/audit_output.txt`).

Two-stage reproducibility caveat. The two-stage pipeline requires per-patch embeddings (`results/probe/patch_embeddings.pkl`, 8.2 GB, not committed; regenerable on GPU via `scripts/extract_patch_embeddings.py`) and hard-codes a Linux working directory. The committed evidence is the aggregate `results/two_stage_results.json` (agreement-based, see App. D); the per-example decisions are not committed.

SITE secondary-subset audit note (corrected before submission). An earlier derived analysis (`compare_site_zeroshot_lora.py` v1, which fed the initial `zeroshot_image_metrics.json`-adjacent figures and the first canonical site table) reconstructed the secondary subset by applying `ORIENTATION_KEYWORDS` to the *question text only*, yielding $n=1,824$ (47.3% raw / 22.6% CAA). The frozen protocol’s pre-specified definition — `results/site/site_protocol.json` → `frozen_ids.secondary` (3,272 examples; 2,024 with images), intersected with the 2,591 evaluated IDs — is the authoritative subset, and all numbers in this paper use it ($n=2,024$; zero-shot 48.3% / 24.0% CAA; VSR-LoRA 48.7% / 24.5%, $p=0.70$). The derived question-only figure was corrected before submission; both definitions support the same conclusion (orientation $\ll$ primary), but only the frozen-ID set matches the pre-specification. No prediction CSV was modified and no model was rerun.

Data licenses/access. VSR: `cambridgelt1/vsr_random` (HF). SITE: `franky-veteran/SITE-Bench`, CC-BY-4.0 (media in 130 GB zips; requires HF read token). LoRA adapters are committed and loadable with `peft`; base-model weights via HF Hub.

Exact artifacts for the probe tables (App. C): `results/probe/probe_results.json`, `results/probe/patch_probe_results.json`, `results/probe/grounded_probe_results.json`, `results/probe/clear_subset_results.json`, `results/probe/grounded_boxes.json`.



#145 The bench is over the cat.

Ground truth: FALSE | Model: TRUE | raw: True.
over - over relation: 44.0% overall (worst, n>=10)



#210 The pizza is right of the bottle.

Ground truth: TRUE | Model: FALSE | raw: False.
right of - right of: 55.2% - directional failure



#157 The chair is left of the teddy bear.

Ground truth: TRUE | Model: FALSE | raw: False.
left of - left of misses here



#33 The cat is facing away from the sink.

Ground truth: TRUE | Model: FALSE | raw: False.
facing away from - facing away from: 53.8% - orientation



#39 The oven is in front of the cake.

Ground truth: TRUE | Model: FALSE | raw: False.



#10 The motorcycle is in front of the truck.

Ground truth: FALSE | Model: TRUE | raw: True.

Figure 5: Representative persistent failures (annotated grid; 15 cases). Images: VSR test set. Per the single-annotator exploratory audit, the most frequent class is *clear-image reasoning failure*: the pose is visually interpretable to a human and the model still errs.